Larch Documentation

Release 0.9.12

Matthew Newville

July 06, 2012

CONTENTS

Larch is a scientific data processing language that is designed to be easy to use for novices while also being complete enough complete enough for advanced data processing and analysis Larch provides a wide range of functionality for dealing with arrays of scientific data, and basic tools to make it easy to use and organize complex data. Written in Python, and making heavy use of the wonderful libraries numpy, scipy, h5py, and matplotlib, Larch has syntax very close to Python, and can be easily extended with Python.

Larch has been primarily developed for dealing with x-ray spectroscopic and scattering data, especially the kind of data collected at modern synchrotrons and x-ray sources. It has several related target application areas:

- XAFS analysis, becoming version Ifeffit Version2.

- tools for micro-XRF mapping visualization and analysis.

- quantitative XRF analysis.

- X-ray standing waves and surface scattering analysis.

- Data collection software for the above.

The idea is that having these different problem areas connected by a common *macro language* will strengthen the analytic tools for all of them.

ONE

LARCH: MOTIVATION AND OVERVIEW

Larch is a scientific data processing language. Though fairly general purpose, it has been developed for and aimed especially at the problems of processing and analyziing x-ray spectroscopic and scattering data collected at modern synchrotrons and x-ray sources. Thus, Larch has several related target application areas, including XAFS, XRF, and X-ray standing waves. The initial movitation was to replace the Ifeffit package for XAFS analysis, but the ability to add XRF capabilites and to use it as a "Spec-like" language for data collection and processing were quickly added to the list of things Larch should be able to do.

Larch gives a simple command-line interface (a Read-Eval-Print Loop), but can also be scripted in "batch mode". GUIs can be built upon Larch by simply generating the commands, making it easy to separate the layout from the actual processing steps, so that the processing steps might be recorded and used to make a "batch script". Larch is easily called from Python. In addition, you can use Larch with remote-procedure calls, so that it can be run from different languages, or run on different machines.

1.1 General Motivation of the Design

Many scientific data collection, visualization, and analysis programs have a *macro language* built into them. These *macro languages* often have built-in commands and datastructures important for the problems being solved. They typically allow customization, automation, scripting, and extension of the fundamental operations needed to get the data collection, visualization, and analysis work done. Some analysis programs have full-fledged languages or are simply built on top of a framework such as Matlab, IDL, Mathematica, R, Eclipse, or Emacs, while many other programs have very simplistic (often buggy) and limited languages.

While *Domain Specific Languages* (the term *macro language* often implies that they are implemented by string substitutions, which may sometimes be true, but is sort of beside the point here) can be a very efficient way to provide flexible interaction and customization of complex software, there are a great many of them in use, making communication and sharing data between programs very hard.

Larch is an attempt to make a domain specific language that can be the basis for x-ray data collection and analysis programs, so that the algorithms and techniques for visualization and analysis can be better shared between different programs and fields. In this respect, Larch is meant to be the foundation or framework upon which data collection, visualization, and analysis programs can be written. By having a common, extensible language and analysis environment, the hope is that it will be easier to make data collection, visualization, and analysis programs interact.

1.2 Overview

Larch is written in Python, an interpreted (non-compiled) language that has become quite popular in a range of scientific disciplines. Python is known for its very clear syntax and readability, and Larch has syntax that is built upon, and very closely related to Python's. Using Python in this way not only gives a fairly elegant and readable language, but

also allows Larch to build upon many great efforts in Python, especially for scientific computing, including numpy, scipy, h5py, and matplotlib. Using Python also turns out to make implementing Larch and adding complex functionality such as XAFS analysis capabilities simple.

In fact, Larch is so closely related to Python that a few key points should be made:

1. All data items in Larch are really Python objects.

2. All Larch code is translated into Python and then run using builtin Python tools.

In a sense, Larch is a dialect of Python. This means has a few that an understanding Larch and Python are close to one another. This in itself can be seen as an advantage – Python is a popular, open-source, language that any programmer can easily learn, and books and web documentation about Python are plentiful. If you known Python, Larch will be very easy to use, and vice versa.

1.3 Design Principles, Key Concepts

Since Larch is intended for processing scientific data, organization of data is a key consideration. The main feature that Larch uses to help the user with organizing data is deceptively simple and useful – the **Group**. This is simply an empty container into which any sort of data can be placed, including other Groups. This provides a heirarchical structure of data that can be accessed and manipulated easily via attributes, as with:

```
larch> my_group = group(x = 0.1*arange(101), title='group 1')
larch> my_group.y = sqrt(my_group.x)
larch> plot(my_group.x, my_group.y, title=my_group.title)
```

That is to say, the Group 'my_group' holds data in a convenient namespace. You can see thed contents of a group:

```
larch> show(my_group)
== Group 0x6e15970: 3 symbols ==
  title: 'group 1'
  x: array<shape=(101,), type=dtype('float64')>
  y: array<shape=(101,), type=dtype('float64')>
```

which shows that this group has 3 components. Other things to note are that 'x' and 'y' hold arrays of data, and thatfunctions such as 'sqrt' act on all elements of the array at once.

Since much of what Larch is used for is modeling or fitting small data sets, another key organizing principle is the **Parameter**. This holds a value that you might want to be optimized in a least-squares fit. Thus, a Parameter can be flagged as a variable, or fixed to not be varied. In addition, it can be given a mathematical expression in terms of other Parameters to determine its value as a constrained value.

1.4 Capabilities

At this writing, Larch has the following general capabilities:

- a full suite of mathematical functionality, with array handling builtin (so that functions work on full arrays).

- a general purpose language with flow-control (for and while loops), and conditional evaluation (if-then-else).

- some built-in I/O functionality for ASCII files and HDF5.

- simple line plots, with customizable line properties.

- simple 2-D image dispays, with some rudimentary customization.

- general-purpose minimization and curve-fitting.

For XAFS analysis in particular, Larch is able to do most data processing and analysis steps needed, including:

- pre-edge background subtraction and normalization
- background subtraction for isolating chi(k)
- XAFS Fourier transforms
- reading and manipulating Feff Path files
- fitting Feff Paths to XAFS data
- general-purpose minimization and curve-fitting.

DOWNLOADING AND INSTALLATION

2.1 Prerequisites

Larch requires Python version 2.6 or higher. Support for Python 3.X is partial, in that the core of Larch does work, and numpy, and scipy, and matplotlib have all been ported to Python 3.X. But the testing for Python 3.X has been minimal, and the graphical interfaces, based on wxWidgets, has not yet been ported to Python 3.X.

Numpy, matplotlib, and wxPython are all required for Larch, and scipy is strongly encouraged (and some functionality depends on it). These are simply installed as standard packages on almost all platforms.

All development is being done through the larch github repository. To get a read-only copy of the atest version, use:

```
git clone http://github.com/xraypy/xraylarch.git
```

2.2 Installation

Installation from source on any platform is:

```
python setup.py install
```

We'll build and distribute Windows binaries and use the Python Package Index soon....

2.3 License

This code and all material associated with it are distributed under the BSD License:

```
Copyright, Licensing, and Re-distribution
-----------------------------------------

Unless otherwise stated, all files included here are copyrighted and
distributed under the following license:

  Copyright (c) 2010-2011 Matthew Newville, The University of Chicago

  Redistribution and use in source and binary forms, with or without
  modification, are permitted provided that the following conditions are met:

  * Redistributions of source code must retain the above copyright notice,
    this list of conditions and the following disclaimer.
  * Redistributions in binary form must reproduce the above copyright
```

notice, this list of conditions and the following disclaimer in the
documentation and/or other materials provided with the distribution.
* The names of the copyright holders or contributers may not be used to
endorse or promote products derived from this software without specific
prior written permission.

THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS
IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO,
THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR
PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT HOLDER OR
CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL,
EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO,
PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR
PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF
LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING
NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS
SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

LARCH TUTORIAL

At its core, Larch is a language for processing of scientific data. This chapter describes the Larch language and introduces data processing using Larch. An important goal of Larch is to make writing and modifying data analysis scripts as simple as possible.

Although aimed at the novice programmer, this tutorial does make a few assumptions about the readers experience with scientific programming. For example, the reader is expected to have a technical background and some familiarity with using scientific data analysis programs. Some understanding of the concepts of how scientific data is stored on computers and of the basics of programming will greatly help the reader.

The Larch language is implemented in Python, and heavily based upon it. Knowledge of Python will greatly simplify learning Larch, and vice versa. This shared syntax is intentional, so that as you learn Larch, you will also be learning Python, which can be used to extend Larch. Alternatively, knowledge of Python will make Larch easy to learn. For further details on Python, including tutorials, see the Python documentation at http://python.org/

Contents:

3.1 Tutorial: Getting Started

This tutorial expects that you already have Larch installed and can run either the program larch, basic Larch interpreter, or larch_gui, the enhanced GUI interpreter:

```
C:> larch
  Larch 0.9.7  M. Newville, T. Trainor (2011)
  using python 2.6.5, numpy 1.5.1
larch>
```

For Windows and Mac OS X users, executable applications will be available.

3.1.1 Larch as a Basic Calculator

To start with, Larch can be used as a simple calculator:

```
larch> 1 + 2
3
larch> sqrt(4.e5)
632.45553203367592
larch> sin(pi/3)
0.8660254037844386
```

You can create your own variables holding values, by assigning names to values, and then use these in calculations:

```
larch> hc = 12398.419
larch> d = 3.13556
larch> energy = (hc/(2*d)) / sin(10.0*pi/180)
larch> print energy
11385.470119348252
larch> angle = asin(hc/(10000*2*d))*180/pi
larch> print angle
11.402879992850263
```

Note that parentheses are used to group multiplication and division, and also to hold the arguments to functions like
`sin()`.

Variable names must start with a letter or underscore ('_'), followed by any number of letters, underscores, or numbers.
You may notice that a dot ('.') may appear to be in many variable names. This indicates an *attribute* of a variable –
we'll get to this in a later section.

If you're familiar with other programming langauges, an important point for Larch (owing to its Python origins) is that
variables are created *dynamically*, they are not pre-defined to have some particular data type. In fact, the a variable
name (say, 'angle' above) can hold any type of data, and its type can be changed easily:

```
larch> angle = 'now I am a string'
```

Although the types of values for a variable can be changed dynamically, values in Larch (Python) do have a definite
and clear type, and conversion between types is rigidly defined – you can add an integer and a real number to give a
real number, but you cannot add a string and a real number. In fact, writing:

```
larch> angle = asin(hc/(10000*2*d))*180/pi
```

is usually described as "create a variable 'angle' and set its value to the result calculated (11.4...)". For those used
to working in C or Fortran, in which variables are static and must be of a pre-declared type, this description is a bit
misleading. A better way to think about it is that the calculation on the right-hand-side of the equal sign results in a
value (11.4...) and we've assigned the name 'angle' to hold this value. The distinction may seem subtle, but it can
have some profound results, as we'll see in the following section when discussing lists and other dynamic values.

3.2 Tutorial: Basic and Complex Data Types

This section of the Larch tutorial describes the types of data that Larch has, and how to use them.

3.2.1 Basic Data Types

As with most programming languages, Larch has several built-in data types to express different kinds of data. These
include the usual integers, floating point numbers, and strings common to many programming languages. A variable
name can hold any of these types (or any of the other more complex types we'll get to later), and does not need to be
declared beforehand or to change its value or type. Some examples:

```
larch> a = 2
larch> b = 2.50
```

The normal '+', '-', '*', and '/' operations work on numerical values for addition, subtraction, multiplication, and
division. Exponentiation is signified by '**', and modulus by '%'. Larch uses the '/' symbol for division or 'true
division', giving a floating point value if needed, even if the numerator and denominator are integers, and '//' for
integer or 'floor' division. Thus:

```
larch> 3 + a
5
```

```
larch> b*2
5.0
larch> 3/2
1.5
larch> 3//2
1
larch> 7 % 3
1
```

Several other operators are supported for bit manipulation.

Literal strings are created with either matched closing single or double quotes:

```
larch> s = 'a string'
larch> s2 = "a different string"
```

A string can include a 'n' character (for newline) or 't' (for tab) and several other control characters as in many languages. For strings that may span over more than 1 line, a special "triple quoted" syntax is supported, so that:

```
larch> long_string = """Now is the time for all good men
.....> to come to the aid of their party"""
larch> print long_string
Now is the time for all good men
to come to the aid of their party
```

It is important to keep in mind that mixing data types in a calculation may or may not make sense to Larch. For example, a string cannot be added to a integer:

```
larch> '1' + 1
Runtime Error:
cannot concatenate 'str' and 'int' objects
<StdInput>
     '1' + 1
```

but you can add an integer and a float:

```
larch> 1 + 2.5
3.5
```

and you can multiply a string by an integer:

```
larch> 'string' * 2
'stringstring'
```

Larch has special variables for boolean or logical operations: `True` and `False`. These are equal to 1 and 0, respectively, but are mostly used in logical operations, which include operators 'and', 'or', and 'not', as well as comparison operators '>', '>=', '<', '<=', '==', '!=', and 'is'. Note that 'is' expresses identity, which is a slightly stricter test than '==' (equality), and is most useful for complex objects.:

```
larch> 2 > 3
False
larch> (b > 0) and (b <= 10)
True
```

The special value `None` is used as a null value throughout Larch and Python.

Finally, Larch knows about complex numbers, using a 'j' to indicate the imaginary part of the number:

```
larch> sqrt(-1)
Warning: invalid value encountered in sqrt
nan
```

```
larch> sqrt(-1+0j)
1j
larch> 1j*1j
(-1+0j)
larch> x = sin(1+1j)
larch> print x
(1.2984575814159773+0.63496391478473613j)
larch> print x.imag
0.63496391478473613
```

To be clear, all these primitive data types in Larch are derived from the corresponding Python objects, so you can consult python documentation for further details and notes.

3.2.2 Objects and Groups

Since Larch is built upon Python, an object-oriented programming language, all named quantities or **variables** in Larch are python objects. Because of this, most Larch variables come with built-in functionality derived from their python objects. Though Larch does not provide a way for the user to define their own new objects, this can be done with the Python interface.

Objects

All Larch variables are Python objects, and so have a well-defined **type** and a set of **attributes** and **methods** that go with it. To see the Python type of any variable, use the builtin `type()` function:

```
larch> type(1)
<type 'int'>
larch> type('1')
<type 'str'>
larch> type(1.0)
<type 'float'>
larch> type(1+0j)
<type 'complex'>
larch> type(sin)
<type 'numpy.ufunc'>
```

The attributes and methods differ for each type of object, but are all accessed the same way – with a '.' (dot) separating the variable name or value from the name of the attribute or method. As above, complex data have `real` and `imag` attributes for the real and imaginary parts, which can be accessed:

```
larch> x = sin(1+1j)
larch> print x
(1.2984575814159773+0.63496391478473613j)
larch> print x.imag
0.63496391478473613
```

Methods are functions that belong to an object (and so know about the data in that object). They are also objects themselves (and so have attributes and methods), but can be called using parentheses '()', possibly with arguments inside the parentheses to change the methods behavior. For example, a complex number has a `conjugate()` method:

```
larch> x.conjugate
<built-in method conjugate of complex object at 0x178e54b8>
larch> x.conjugate()
(1.2984575814159773-0.63496391478473613j)
```

Strings and other data types have many more attributes and methods, as we'll see below.

To get a listing of all the attributes and methods of a object, use the builtin `dir()` function:

```
larch> dir(1)
['__abs__', '__add__', '__and__', '__class__', '__cmp__', '__coerce__', '__delattr__', '__div__', '_
larch> dir('a string')
['__add__', '__class__', '__contains__', '__delattr__', '__doc__', '__eq__', '__format__', '__ge__',
```

Again, we'll see properties of objects below, as we look into more interesting data types, or you can look into Python documentation.

Groups

In addition to using basic Python objects, Larch organizes data into Groups. A Group is simply a named container for variables of any kind, including other Groups. In this way, Groups have a heirarchical structure, much like a directory of files. Each Larch variable belongs to a Group, and can be accessed by its full Group name. The top-level Group is called '_main'. You'll rarely need to use that, but it's there:

```
larch> myvar = 22.13
larch> print _main.myvar
22.13
larch> print myvar
22.13
```

You can create your own groups and add data to it with the builtin `group()` function:

```
larch> g = group()
larch> g
<Group>
```

You can add variables to your Group 'g', using the '.' (dot) to separate the parent group from the child object:

```
larch> g.x = 1002.8
larch> g.name = 'here is a string'
larch> g.data = arange(100)
larch> print g.x/5
200.56
```

(`arange()` is a builtin function to create an array of numbers). As from the above discussion of objects, the '.' (dot) notation implies that 'x', 'name', and 'data' are attributes of 'g' – that's entirely correct. Groups have no other properties than the data attributes (and functions) you add to them. Since they're objects, you can use the `dir()` function as above:

```
larch> dir(g)
['data', 'name', 'x']
```

(Note that the order shown may vary). You can also use the builtin `show()` function to get a slightly more complete view of the group's contents:

```
larch> show(g)
== Group: 3 symbols ==
  data: array<shape=(100,), type=dtype('int32')>
  name: 'here is a string'
  x: 1002.8
```

The `group()` function can take arguments of attribute names and values, so that this group could have been created with a single call:

```
larch> g = group(x=1002.8, name='here is a string', data=arange(100))
```

Many Larch functions will return groups or take a 'group' argument to write data into. That is, a function that reads data from a file will almost certainly organize that data into a group, and simply return the group for you to name, perhaps something like:

```
larch> cu = read_ascii('cu_150k.xmu')
```

Builtin Larch Groups

Larch starts up with several groups, organizing builtin functionality into different groups. The top-level '_main' group begins with 3 principle subgroups, '_builtin', '_sys', and '_math' for basic functionality. For almost all uses of Larch, several additional groups are created for more specific functionality are created on startup by Larch plugins. The principle starting groups are describe in *Table of Basic Larch Groups*

Table of Basic Larch Groups. These groups are listed in order of how they will be searched for functions and data.

Group Name	description
_builtin	basic builtin functions.
_math	mathematical and array functions.
_sys	larch sstem-wide variables.
_io	file input/output functions.
_plotter	plotting and image display functions.
_xafs	XAFS-specific functions.

The functions in '_builtin' are mostly inherited from Python's own built-in functions. The functions in '_math' are mostly inherited from Numpy, and contain basic array handling and math.

How Larch finds variable names

With several builtin groups, and even more groups created to store your own data to be processed, Larch ends up with a complex heirarchy of data. This gives a good way of organizing data, but it also leads to a question of how variable names are found. Of course, you can always access a function or data object by its full name:

```
larch> print _math.sin(_math.pi/2)
1.0
```

but that's too painful to use, and of course, one needs to be able to do:

```
larch> print sin(pi/2)
1.0
```

and have Larch know that when you say `sin()`, you mean `_math.sin()`. The way this look-up of names works is that Larch keeps a list of groups that it will search through for names. This list is held in the variable `_sys.searchGroups`, and can be viewed and modified during a Lach session. On startup, this list has the groups listed in *Table of Basic Larch Groups*, in the order shown. To be clear, if there was a variable named `_sys.something` and a `_math.something`, typing 'something' would resolve to `_sys.something`, and to access `_math.something` you would have to give the full name. For the builtin functions and variables, such clashes are not so likely, but they are likely if you read in many data sets as groups, and want to access the contents of the different groups.

3.2.3 More Complex Data Structures: Lists, Arrays, Dictionaries

Larch has many more data types built on top of the primitive types above. These are generally useful for storing collections of data, and can be built up to construct very complex structures. These are all described in some detail

here. But as these are all closely related to Python objects, further details can be found in the standard Python documentation.

Lists

A list is an ordered sequence of other data types. They are **heterogeneous** – they can be made up of data with different types. A list is constructed using brackets, with commas to separate the individual:

```
larch> my_list1 = [1, 2, 3]
larch> my_list2 = [1, 'string', sqrt(7)]
```

A list can contain a list as one of its elements:

```
larch> nested_list = ['a', 'b', ['c', 'd', ['e', 'f', 'g']]]
```

You can access the elements of a list using brackets and the integer index (starting from 0):

```
larch> print my_list2[1]
'string'
larch> print nested_list[2]
['c', 'd', ['e', 'f', 'g']]
larch> print nested_list[2][0]
'c'
```

Lists are **mutable** – they can be changed, in place. To do this, you can replace an element in a list:

```
larch> my_list1[0] = 'hello'
larch> my_list1
['hello', 2, 3]
```

As above, lists are python **objects**, and so come with methods for interacting with them. For example, you can also change a list by appending to it with the 'append' method:

```
larch> my_list1.append('number 4, the larch')
larch> my_list1
['hello', 2, 3, 'number 4, the larch']
```

All lists will have an 'append' method, as well as several others:

- count – to return the number of times a particular element occurss in the list
- extend – to extend a list with another list
- index – to find the first occurance of an element
- insert – to insert an element in a particular place.
- pop – to remove and return the last element (or other specified index).
- remove – remove a particular element
- reverse – reverse the order of elements
- sort – sort the elements.

Note that the methods that change the list do so *IN PLACE* and return `None`. That is, to sort a list, do this:

```
larch> my_list.sort()
```

but not this:

```
larch> my_list = my_list.sort()    # WRONG!!
```

as that will set 'my_list' to None.

You can get the length of a list with the built-in `len()` function, and test whether a particular element is in a list with the *in* operator:

```
larch> my_list = ['a', 'b', 'c', 'd', 'e', 'f', 'g', 'h', 'i', 'j']
larch> print len(my_list)
10
larch> 'e' in my_list
True
```

You can access a sub-selection of elements with a **slice**, giving starting and ending indices between brackets, separated by a colon. Of course, the counting for a slice starts at 0. It also excludesthe final index:

```
larch> my_list[1:3]
['b', 'c']
larch> my_list[:4]    # Note implied 0!
['a', 'b', 'c', 'd']
```

You can count backwards, and using '-1' is a convenient way to get the last element of a list. You can also add an optional third value to the slice for a step:

```
larch> my_list[-1]
'j'
larch> my_list[-3:]
['h', 'i', 'j']
larch> my_list[::2]   # every other element, starting at 0
['a', 'c', 'e', 'g', 'i']
larch> my_list[1::2]   # every other element, starting at 1
['b', 'd', 'f', 'h', 'j']
```

A final important property of lists, and of basic variable creation in Larch (and Python) is related to the discussion above about variable creation and assignment. There we said that 'creating a variable':

```
larch> my_list = ['a', 'b', 'c', 'd', 'e', 'f', 'g', 'h', 'i', 'j']
```

was best thought of as creating a value (here, the literal list "['a', 'b', ..., 'j']") and then assigning the name 'my_list' to point to that value. Here's why we make that distinction. If you now say:

```
larch> your_list = my_list
```

the variable 'your_list' now points to the same value – the same list. That is, it does not make a copy of the list. Since the list is mutable, changing 'your_list' will also change 'my_list':

```
larch> your_list[0] = 500
larch> print my_list[:3]
[500, 'b', 'c']                  # changed!!
```

You can make a copy of a list, by selecting a full slice:

```
larch> your_list = my_list[:]
larch> your_list[0] = 3.2444
larch> print my_list[:3]
[500, 'b', 'c']                  # now unchanged

larch> your_list[0] == my_list[0]
False
```

Note that this behavior doesn't happen for immutable data types, including the more primitive data types such as integers, floats and strings. This is essentially because you cannot assign to parts of those data types, only set its entire value.

As always, consult the Python documentation for more details.

Tuples

Like lists, tuples are sequences of heterogenous objects. The principle difference is that tuples are **immutable** – they cannot be changed once they are created. Instead, tuples are a simple ordered container of data. The syntax for tuples uses comma separated values inside (optional!) parentheses in place of brackets:

```
larch> my_tuple = (1, 'H', 'hydrogen')
```

Like lists, tuples can be indexed and sliced:

```
larch> my_tuple[:2]
(1, 'H')
larch> my_tuple[-1]
'hydrogen'
```

Due to their immutability, tuples have only a few methods ('count' and 'index' with similar functionality as for list).

Though tuples they may seem less powerful than lists, and they are actually used widely with Larch and Python. In addition to the example above using a tuple for a short, fixed data structure, many functions will return a tuple of values. For this case, the simplicity an immutability of tuples is a strength becaues, once created, a tuple has a predictable size and order to its elements, which is not true for lists. That is, if a larch procedure (which we'll see more of below) returns two values as a tuple:

```
larch> def sumdiff(x, y):
.....>     return x+y, x-y
.....> enddef
larch> x = sumdiff(3, 2)
larch> print x[0], x[1]
5 1
```

Because the returned tuple has a fixed structure, you can also assign the it directly to a set of (the correct number of) variables:

```
larch> plus, minus = sumdiff(10, 3)
larch> print plus, minus
13 7
```

A second look at Strings

Though discussed earlier in the basic data types, strings are closely related to lists as well – they are best thought of as a sequence of characters. Like tuples, strings are actually immutable, in that you cannot change part of a string, instead you must create a new string. Strings can be indexed and sliced as with lists and tuples:

```
larch> name = 'Montaigne'
larch> name[:4]
'Mont'
```

Strings have many methods – over 30 of them, in fact. To convert a string to lower case, use its `lower()` method, and so on:

```
larch> 'Here is a String'.lower()
'here is a string'
larch> 'Here is a String'.upper()
'HERE IS A STRING'
larch> 'Here is a String'.title()
'Here Is A String'
```

This aslo shows that the methods are associated with strings themselves – even literal strings, and simply with variable names.

Strings can be split into words with the `split()` method, which splits a string on whitespace by default, but can take an argument to change the character (or substring) to use to split the string:

```
larch> 'Here is a String'.split()
['Here', 'is', 'a', 'String']

larch> 'Here is a String'.split('i')
['Here ', 's a Str', 'ng']
```

As above, this is really only touching the tip of the iceberg of string functionality, and consulting standard Python documentation is recommended for more information.

Arrays

Whereas lists are sequences of heterogeneous objects that can grow and shrink, and included deeply nested structures, they are not well suited for holding numerical data. Arrays are sequences of the same primitive data type, and so are much closer to arrays in C or Fortran. This makes them much more suitable for numeric calculations, and so are extremely important in Larch. There are many ways to create arrays, including the builtin `array()` function which will attempt to convert a list or tuple of numbers into an Array. You can also use the builtin `arange()` function to create an ordered sequence of indices ([1, 2, 3, ...]), or one of several other methods to create arrays.

Arrays are so important for processing numerical data that the next section is devoted to them.

Dictionaries

Our final basic data-structure is the dictionary, which is a container that maps values to keys. This is sometimes called a hash or associative array. Like a list, a dictionary holds many heterogeneous values, and can be altered in place. Unlike a list, the elements of a dictionary have no guaranteed order, and are not selected by integer index, and multiple values cannot be selected by a slice. Instead, the elements of a dictionary are accessed by key, which is normally a string, but can also be an integer or floating point number, or even a tuple or some other objects – any **immutable** object can be used. Dictionaries are delimited by curly braces, with colons (':') separating key and value, and commas separating different elements:

```
larch> atomic_weight = {'H': 1.008, 'He': 4.0026, 'Li': 6.9, 'Be': 9.012}
larch> print atomic_weight['He']
4.0026
```

You can also add more elements to a dictionary by assigning to a new key:

```
larch> atomic_weight['B']  = 10.811
larch> atomic_weight['C']  = 12.01
```

Dictionaries have several methods, such as to return all the keys or all the values, with:

```
larch> atomic_weight.keys()
['Be', 'C', 'B', 'H', 'Li', 'He']
```

```
larch> atomic_weight.values()
[9.0120000000000005, 12.01, 10.811, 1.008, 6.9000000000000004, 4.0026000000000002]
```

Note that the keys and values are not in the order they were entered in, but do have the same order.

As with lists, dictionaries are mutable, and the values in a dictionary can be any object, including other lists and dictionaries, so that a dictionary can end up with a very complex structure. Dictionaries are quite useful, and are in fact used throughout python.

3.3 Tutorial: Basic Array Manipulation

3.3.1 Functions for creating arrays

Table of Array Creation Functions

Table of Array Creation Functions. These functions can all be used to create arrays in Larch.

Function Name	description	example
array	array from list	arr = array([1,2,3])
arange	indices 0, N-1	arr = arange(10)
zeros	fill with N zeros	arr = zeros(10)
ones	fill with N ones	arr = ones(10)
linspace	fill with bounds and N	arra = linspace(0, 1, 11)

Some examples of using these functions are needed:

```
larch> i = arange(10)
larch> i
array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
larch> f = arange(10, dtype='f8')
larch> f
array([0., 1., 2., 3., 4., 5., 6., 7., 8., 9.])
larch> c = arange(10, dtype='c16')
larch> c
array([ 0.+0.j,  1.+0.j,  2.+0.j,  3.+0.j,  4.+0.j,  5.+0.j,  6.+0.j,
        7.+0.j,  8.+0.j,  9.+0.j])
```

Here, the **dtype** argument sets the data type for the array members – in this case 'f8' means '8 byte floating point' and 'c16' means '16 byte complex' (i.e, double precision, and double precision complex, respectively).

The linspace() function is particularly useful for creating arrays, as it takes a starting value, ending value, and number of points between these:

```
larch> s = linspace(0, 10, 21)
larch> s
array([  0. ,   0.5,   1. ,   1.5,   2. ,   2.5,   3. ,   3.5,   4. ,
         4.5,   5. ,   5.5,   6. ,   6.5,   7. ,   7.5,   8. ,   8.5,
         9. ,   9.5,  10. ])
```

Several variants are possible. For more information, consult the numpy tutorials, or use the online help system within Larch (which will print out the documentation string from the underlying numpy function):

```
larch> help(linspace)

    Return evenly spaced numbers over a specified interval.

    Returns `num` evenly spaced samples, calculated over the
```

```
interval ['start', 'stop' ].

The endpoint of the interval can optionally be excluded.

Parameters
----------
start : scalar
    The starting value of the sequence.
stop : scalar
    The end value of the sequence, unless `endpoint` is set to False.
    In that case, the sequence consists of all but the last of ``num + 1``
    evenly spaced samples, so that `stop` is excluded.  Note that the step
    size changes when `endpoint` is False.
num : int, optional
    Number of samples to generate. Default is 50.
endpoint : bool, optional
    If True, `stop` is the last sample. Otherwise, it is not included.
    Default is True.
retstep : bool, optional
    If True, return (`samples`, `step`), where `step` is the spacing
    between samples.

....
```

3.3.2 The array *dtype* – datatype

3.3.3 Slicing and extracting sub-arrays

3.4 Tutorial: Conditional Execution and Flow Control

Two important needs for a full-featured language are the ability to run different statements under different conditions, and to repeat certain calculations. These are generally called 'flow control', as these statements control how the program will flow through the text of the script. In the discussion here, we will also introduce a few new concepts and Larch statements.

So far in this tutorial, all the text written to the Larch command line has been a single line of text that is immediately run, either printing output to the terminal or assigning a value to a variable. These are both examples of **statements**, which are the basic pieces pf text you send to the program. So far we've seen three types of statements:

1. simple statements or expressions, such as:

   ```
   larch> 1+sqrt(3)
   ```

   or:

   ```
   larch> atomic_weight.keys()
   ```

   where we have an expression evaluated. At the command line, these values are printed – in a script they would not be printed.

2. print statements, such as:

   ```
   larch> print sqrt(3)
   ```

   where we explicitly command larch to print the evaluated expression – this would print if run from a script.

3. assignment statements, such as:

```
larch> x = sqrt(3)
```

where we assign the name 'x' to hold the evaluated expression.

In fact, though these are the most common types of statements, there are many more types of statements. We will introduce a few more statement types here, including compound statements that take up more than one line of test.

3.4.1 Conditional Evaluation with if and else

A fundamental need is to be able to execute some statement(s) when some condition is met. This is done with the **if** statement, an example of which is:

```
larch> if x == 0: print 'x is 0!'
```

Which will print 'x is 0!' if the value of *x* is equal to 0. The *x == 0* in this if statement is called the **test**. A test is a Boolean mathematical expression. While most usual mathematical expressions use operators such as '+', '-', a Boolean expression uses the operators listed in *Table of Boolean Operators* to evaluate to a value of `True` or `False`.

A single-line if statement as above looks like this:: if <test>: statement

The 'if' and ':' are important, while '<test>' can be any Boolean expression. If the test evaluates to `True`, the statement is executed.

If statements can execute multiple statements by putting the statements into a "block of code":

```
if x == 0:
    print 'x is equal to 0!'
    x = 1
endif
```

Which is to say that the multiple-line form of the if statement looks like this:

```
if <test>:
  <statements>
endif
```

where '<statements>' here means a list of statements, and the 'endif' is required (see *Code Block Ends*). For the above, two statements will be run if x is equal to 0 – there is no restriction on how many statements can be run.

An 'else' statement can be added to execute code if the test is False:

```
if x == 0:
    print 'x is equal to 0!'
    x = 1
else
    print 'x is not 0'
endif
```

Multiple tests can be chained together with the 'elif' (a contraction of 'else if'):

```
if x == 0:
    print 'x is equal to 0!'
elif x > 0:
    print 'x is positive'
elif x > -10:
    print 'x is a small negative number'
else:
    print 'x is negative'
endif
```

Here the 'x > 0' test will be executed if the 'x == 0' test fails, and the 'x > -10' test will be tried if that fails. **Table of Boolean Operators** The operators here all take the form *right OP left* where OP is one of the operators below. Note the distinction between '==' and 'is'. The former compares *values* while the latter compares the identity of two objects.

boolean operator	meaning
==	has equal value
!=	has unequal value
>	has greater value
>=	has greater or equal value
<	has smaller value
<=	has smaller or equal value
is	is identical to
not	is not `True`
and	both operands are `True`
or	either operand is `True`

Note that in Larch, as in Python, any value can be used as a test, not just values that are `True` or `False`. As you might expect, for example, the value 0 is treated as `False`. An empty string is also treated as `False`, as is an empty list or dictionary. Most other values are interpreted as `True`.

3.4.2 For loops

It is often necessary to repeat a calculation multiple times. A common method of doing this is to use a **loop**, including using a loop counter to iterates over some set of values. In Larch, this is done with a **for loop**. For those familiar with other languages, a Larch for loop is a bit different from a C for loop or Fortran do loop. A for loop in Larch iterates over an ordered set of values as from a list, tuple, or array, or over the keys from a dictionary. Thus a loop like this:

```
for x in ('a', 'b', 'c'):
    print x
endfor
```

will go through values 'a', 'b', and 'c', assigning each value to *x*, then printing the value of x, which will result in printing out:

```
a
b
c
```

Similar to the the *if* statement above, the for loop has the form:

```
for <varlist> in <sequence>:
    <statements>
endfor
```

Compared to a C for loop or Fortran do loop, the Larch for loop is much more like a *foreach* loop. The common C / Fortran use case of interating over a set of integers can be emulated using the builtin `range()` function which generates a sequence of integers. Thus:

```
for i in range(5):
    print i, i/2.0
endfor
```

will result in:

```
0, 0.0
1, 0.5
2, 1.0
```

```
3, 1.5
4, 2.0
```

Note that the builtin `range()` function generates a sequence of integers, and can take more than 1 argument to indicate a starting value and step. It is important to note that the sequence that is iterated order does not be generated from the `range()` function, but can be any list, array, or Python sequence. Importantly, this includes strings(!) so that:

```
for char in 'hello':  print char
```

will print:

```
h
e
l
l
o
```

This can cause a common sort of error, in that you might expect some variabe to hold a list of string values, but it actually holds a single string. Notice that:

```
filelist = ('file1', 'file2')
for fname in filelist:
    fh = open(fname)
    process_file(fh)
    fh.close()
endfor
```

would act very differently if filelist was changed to 'file1'!

Multiple values can be assigned in each iteration of the for loop. Thus, iterating over a sequence of equal-length tuples, as in:

```
for a, b in (('a', 1), ('b', 2), ('c', 3)):
    print a, b
endfor
```

will print:

```
a, 1
b, 2
c, 3
```

This may seem to be mostly of curious interest, but can be extremely useful especially when dealing with dictionaries or with arrays or lists of equal length. For a dictionary *d*, *d.items()* will return a list of two-element tuples as above of key, value. Thus:

```
mydict = {'a':1, 'b':2, 'c':3, 'd':4}
for key, val in mydict.items():
    print key, val
endfor
```

will print (note that dictionaries do no preserve order, but the (key, val) pairs match:

```
a 1
c 3
b 2
d 4
```

The builtin `zip()` function is similarly useful, turning a sequence of lists or arrays into a sequence of tuples of the corresponding elements of the lists or arrays. Thus:

```
larch> a = range(10)
larch> b = sin(a)
larch> c = cos(a)
larch> print zip(a, b, c)
[(0, 0.0, 1.0), (1, 0.8414709848078965, 0.54030230586813977),
 (2, 0.90929742682568171, -0.41614683654714241), ....]
```

(Note that for arrays or lists of unequal length, `zip()` will return tuples until any of its arguments runs out of elements). Thus a for loop can make use of the `zip()` function to iterate over multiple arrays:

```
larch> a = arange(101)/10.0
larch> print 'X    SIN(X)   SIN(Y)\n================\n'
larch> for a, sval, cval in zip(a, sin(a), cos(a)):
.....>     print '%.3f, %.5f, %.5f' % (a, sval, cval)
.....> endfor
```

will print a table of sine and cosine values.

A final utility of note for loops is `enumerate()` which will return a tuple of (index, value) for a sequence. That is:

```
larch> for i, a in ('a', b', 'c'):
.....>     print i, a
.....> endfor
```

will print:

```
0, a
1, b
2, c
```

It is sometimes useful to jump out of a for loop, or go onto the next value in the sequence. The *break* statement will exit a for loop immediately:

```
for fname in filelist:
    status = get_status(fname)
    if status < 0:
       break
    endif
    more_processing(fname)
endfor
print 'processed up to i = ', i
```

may jump out of the loop before the sequence generated by 'range(10)' is complete. The variable 'i' will have the final used value.

To skipover an iteration of a loop but continue on, use the *continue* statement:

```
for fname in filelist:
    status = get_status(fname)
    if status < 0:
       continue
    endif
    more_processing(fname)
endfor
```

3.4.3 While loops

While a for loop generally walks through a pre-defined set of values, a *while* loop executes as long as some test is `True`. The basic form is:

```
while <test>:
    <statements>
endwhile
```

Here, the test works as for *if* – it is a Boolean expression, evaluated at each iteration of the loop. Generally, the expression will test something that has been changed inside the loop (even if implicitly). The classic while loop increases a counter at each iteration:

```
counter = 0
while counter < 10:
    do_something(counter)
    counter =+ 1
endwhile
```

A while loop is easily turned into an infinite loop, simply by not incrementing the counter. Then again, the above loop would easily be converted into a for loop, as the counter is incremented by a fixed amout at each iteration. A more realistic use would be:

```
n = 1
while n < 100:
    n = (n + 0.1) * n
    print n
endwhile
```

An additional use for a while loop is to use an implicit or external condition, such as time:

```
now = time.time() # return the time in seconds since Unix epoch
while time.time() - now < 15:   # That is 'for 15 seconds'
    do_someting()
endwhile
```

The *break* and *continue* statements also work for while loops, just as they do with for loops. These can be used as ways to exit an other-wise infinite while loop:

```
while True:   # will never exit without break!
    answer = raw_input('guess my favorite color>')
    if answer == 'lime':
        break
    else:
        print 'Nope, try again'
    endif
endwhile
```

3.5 Tutorial: Dealing With Errors

When an error exists in the syntax of your script, or an erro happens when running your script, an *Error* or *Exception* is generated, and the execution of your script is stopped.

3.5.1 Syntax Errors

Syntax errors result from incomplete or ill-formed larch syntax. For example:

```
larch> x = 3 *
SyntaxError
<StdInput>
    x = 3  *
```

This indicates that the Larch interpreter could not understand the meaning of the statement 'x = 3 ', *because it excepts a value after the* ''. Syntax errors are spotted and raised before the interpreter tries to evaluate the expression. That is because Larch first fully parses any statements (or block of statements if you're entering multiple statements or loading a script from a file) into a partially compiled, executable form before executing. Because of this two-step approach (first parse to intermediate form, then execute that intermediate form), syntax errors are sometimes referred to as a parsing errors.

3.5.2 Exceptions

Even if the syntax of your script is correct, the logic might not be. In addition, even if the logic is correct for most cases, it might not be correct for all. For example, certain values might cause an error run time:

```
larch> n = 1
larch> print 4.0 / ( n - 1)
ZeroDivisionError('float division')
<StdInput>
    print 4.0/(n-1)
          ^^^
```

which is saying that you can't divide by 0. This is known as a **Runtime Exception**. It might indicate a programming error, for example that you didn't test if the denominator was 0 before doing the division.

Larch (as inherited from Python) has many different types of exceptions, so that dividing by zero, as above, is detected as a different exception from, say, trying to open a file that doesn't exist:

```
larch> fh = open('foo', 'r')
Error running <built-in function open>
IOError(2, 'No such file or directory')
<StdInput>
    fh = open('foo', 'r')
         ^^^
```

or trying to add an integer and a string together:

```
larch> 4 + 'a'
TypeError("unsupported operand type(s) for +: 'int' and 'str'")
<StdInput>
    4 + 'a'
```

Though they are called exceptions, such problems are fairly common when developing programs or writing scripts. Having a built-in way to test for and handle different kinds of exceptions is an important part of many modern computer languages, and Larch has this capability with its **try** and **except** statements.

3.5.3 Try and Except

The **try** statement will execute a block of code and look for certain types of exceptions. One or more **except** statements can be added to specify blocks of code to execute if the specified exception occurs. As a simple example:

```
try:
    x = a/b
except ZeroDivisionError:
    print 'saw a divide by zero!'
    x = 0
endtry
<more statements>
```

If b is not 0, x is set to the value of a/b. If b is 0, executing *x = a/b* will cause a ZeroDivisionError (as we saw above), so the block with the print statement and setting x to 0 will be executed. In either of these cases (no exception, or a handled exception), execution will continue as normal. If a different problem occurs – an "unhandled exception" – such as the case if *a* holds a string value with *b* holds an integer, then execution will stop and the corresponding exception will be raised.

There can be several **except** statements for each **try** statement, to check for multiple types of problems. These will be checked in order. For example:

```
try:
    x = a/b
except ZeroDivisionError:
    print 'saw a divide by zero!
    x = 0
except TypeError:
    print "a and b are of different types -- can't divide"
endtry
<more statements>
```

It is sometimes useful to run certain code only when a looked-for error has not occurred. For example, it is often agood idea to test when opening a file for IOError (which covers a range of issues such as the file not being found), and only reading that file if it actually opened. For example, to read in a file into a list of lines, the recommended practice is to do:

```
try:
    fh = open(filename, 'r')
except IOError:
    print 'cannot open file %s!' % filename
    datalines = []
else:
    datalines = fh.readlines()
    fh.close()
endtry
<operate on datalines>
```

There is a very large number of exception types built into Larch, all inherited from Python. See the standard Python documentation for more details.

3.5.4 Raising your own exceptions

In certain cases, you may want to cause an exception to occur. This need is most likely to happen when writing your own procedures, and want to ensure that the input arguments can be handled correctly.

To cause an exception, you use the **raise** statement, and you are said to be "raising an exception":

```
larch> raise TypeError("wrong data type")
```

3.6 Tutorial: Writing Procedures

Any moderately complex script will eventually include calcluations that need to be repeated. The preferred way to do this is write your own function or **procedure**, which can be called exactly as the built-in functions. For clarification, here we use the word **procedure** for a function written in Larch, leaving **function** to imply a Python function. In fact, as we will see, there is very little difference in practical use.

Once you're ready to write procedures, you'll almost certainly want to read about running Larch scripts and modules in the next section of this tutorial.

3.6.1 Def statement

To define a procedure, you use the **def** statement, and write a block of code. This looks much like the **if**, **for**, and **while** blocks discussed earlier. A simple example would be:

```
def sayhello():
    print 'hello!'
enddef
```

With this definition, one can then run this procedure as you would run any other built-in function, by writing:

```
larch> sayhello()
hello!
```

Of course, you can write procedures that take input arguments, such as:

```
def safe_sqrt(x):
  if x > 0:
     print sqrt(x)
  else:
     print 'Did you want sqrt(%f) = %f?' % (-x,  sqrt(-x))
  endif
enddef
```

Here *x* will hold whatever value is passed to it, so that:

```
larch> safe_sqrt(4)
2.0
larch> safe_sqrt(-9)
Did you want sqrt(9.000000) = 3.000000?
```

Of course, you will most often want a procedure to return a value. This is done with the **return** statement. A **return** statement can be put anywhere in a procedure definition. When encountered, it will cause the procedure to immediately exist, passing back any indicated value(s). If no **return** statement is given in a procedure, it will return `None` when the procedure has fully executed. An example:

```
def safe_sqrt(x):
  if x > 0:
     return sqrt(x)
  else:
     return sqrt(-x)
  endif
enddef
```

which can now be used as:

```
larch> print safe_sqrt(4)
2.0
larch> x = safe_sqrt(-10)
larch> print x
3.16227766017
```

return can take multiple arguments, separated by a comma, which is to say a **tuple**. As an example:

```
larch> def sum_diff(x, y):
.....>    return x + y, x-y
.....> enddef
larch> print sum_diff(3., 4.)
(7.0, -1.0)
```

This is discussed in more detail below.

The formal definition of a procedure looks like:

```
def <procedure_name>(<arguments>):
    <block of statements>
enddef
```

3.6.2 Namespace and "Scope" inside a Procedure

While inside a procedure, and important consideration is "what variables does this procedure have access to?". Generally speaking, there is no reason to expect it to know about any variables that are not passed in as arguments or created inside the procedure definition.

3.6.3 The return statement, and multiple Return values

As seen above, the **return** statement will exist a procedure, and send back a value to the calling code. The return value can be either a single value or a tuple of values, which gives a convenient way to return multiple values from a single procedure. Thus:

```
larch> def my_divmod(x, y):
.....>    return (x // y, x % y)  # note use of // for integer division!
.....> enddef
larch> print my_divmod(100, 7)
14, 2
```

But be careful when assigning the return value to variable(s). You can do:

```
larch> xdiv, xmod = my_divmod(100, 7)
larch> print xdiv
14
```

or:

```
larch> result = my_divmod(100, 7)
larch> print result[0], result[1]
14, 2
```

Because a return value from a procedure can hold many values, it is best to be careful when writing a procedure that you document what the return value is, and when using a procedure that you're getting the correct number of values.

3.6.4 Keyword arguments

For the procedures defined so far, the arguments have been both required and in a fixed order. Sometimes, you'll want to give a procedure optional arguments, and perhaps allow some flexibility in the order of the arguments. Larch allows this with **keyword** arguments. In a procedure definition, you add an argument name with a default value, like this:

```
def xlog(a, base=e):
    """return log(a) with base = base (default=e=2.71828...)
    """
     if base > 1:
        return log(a) /log(base)
    else:
        print 'cannot calculate log base %f' % base
    endif
enddef
```

Unless passed in, the value of *base* will take the default value of *e*. This can then be used as:

```
larch> xlog(16)
2.7725887222397811
larch> xlog(16, base=10)
1.2041199826559246
larch> xlog(16, base=2)
4.0
```

You can supply many keyword arguments, but they must all come *after* the positional arguments.

A procedure can be written to take an unspecified number of positional and keyword parameters, using a special syntax for unspecified positional arguments and for unspecified keyword arguments. To use unspecified positional arguments, a procedure definition takes an argument preceded by a '*' after all the named positional arguments, like this:

```
def addall(a, b, *args):
    """add all (at least 2!!) arguments given"""
    out = a + b
    for c in args:
        out = out + c
    endfor
    return out
enddef
```

Here, the **'*args'** arguments means to use the variable 'args' to hold any number of positional arguments beyond those explicitly given. Inside the procedure, a tuple named 'args' will hold any positional parameters included in the call to 'addall' past the first two (which will be held by 'a' and 'b'). Thus, this procedure can be used as:

```
larch> addall(2, 3)          # args = ()
5
larch> addall(2, 3, 5, 7)    # args = (5, 7)
17
```

To add support for unspecified keyword parameters, one adds a named argument to the procedure definition preceded by two asterisks: **'**keywords'**. For example:

```
def operate(a, b, **options):
    """perform operation on a and b"""
    debug = options.get('debug', True)
    verbose = options.get('verbose', False)
    op  = options.get('op', 'add')
    if verbose:
       print 'op == %s ' % op
    endif
    if op == 'add':
        return a + b
  elif op == 'sub':
        return a - b
  elif op == 'mul':
        return a * b
  elif op == 'div':
        return a / b
  else:
        if debug:  print 'unsupported operation!'
    endif
enddef
```

As you may have figured out, inside the procedure, 'options' will hold a dictionary of keyword names/values passed into it. With this (perhaps contrived) definition, you can call 'operate' many ways to change its behavior:

```
larch> operate(3, 2, op='add')
5
larch> operate(3, 2, op='add', verbose=True)
op == add
5
larch> operate(3, 2, op='mul', verbose=True)
op == mul
6
larch> operate(3, 2, op='xxx', verbose=True)
op == xxx
unsupported operation!
larch> operate(3, 2, op='xxx', debug=False)
op == xxx
```

As with the **'*args'**, the **'**options'** in the procedure definition must appear after any named keyword parameters, and will not include the named keyword parameters.

3.6.5 Documentation Strings

It is generally a good idea to document your procedures so that you and others can read what it is meant to do and how to use it. Larch has a built-in mechanism for supporting procedure documentaion. If the first statement in a procedure is a **bare string** (that is, a string that is not assigned to a variable), then this will be used as the procedure documentation. You can use triple-quoted strings for multi-line documentation strings. This doc string will be used by the built-in help mechanism, or when viewing details of the procedure. For example:

```
def safe_sqrt(x):
  """safe sqrt function:
 returns sqrt(abs(x))
 """
 return sqrt(abs(x))
enddef
```

With this definition:

```
larch> help(safe_sqrt)
  safe sqrt function:
     returns sqrt(abs(x))
```

3.7 Tutorial: Running Larch Scripts, and Modules

Once you've done any significant amount of work with Larch, you'll want to save what you've done to a file of Larch code, and run it over again, perhaps changing some input. There are a few ways to do this. Writing procedures that can be re-used is a highly recommended approach.

3.7.1 Running a script with run

If you have a file of Larch code, you can run it with the built-in **run** function:

```
# file  myscript.lar
print 'hello from myscript.lar!
for i in range(5):
   print i, sqrt(a)
endfor
#
```

To run this, you simply type:

```
larch> run('myscript.lar')
hello from myscript.lar!
0 0.0
1 1.0
2 1.41421356237
3 1.73205080757
4 2.0
```

A script can contain any Larch code, including procedure definitions. After running the script, any variables assigned in the script will exist in your larch session. For example, after the loop in *myscript.lar*, the variable *i* will be 4, and you can access this variable:

```
larch> print i
4
```

This is to say that the script runs in the "top-level namespace", about which we'll see more below.

3.7.2 Importing a Larch Module

A larch script can also be **imported** using the **import** statement:

```
larch> import myscript
```

Notice a few differences. First, the '.lar' suffix was removed. Second, the name is **NOT IN QUOTES** as one my expect for a string containing a file name. This is because the **import** statement knows what extensions to look for:

- module lookup

- import as..

- from xx import yy as zz

3.7.3 Namespaces, again

3.8 Tutorial: Builtin Functions

3.9 Tutorial: Reading and Writing Data

Larch has several built-in functions for reading scientific data. The intention that the types of supported files will increase. In addition, many Python modules for reading standard types of image data can be used.

3.9.1 Simple ASCII Column Files

A simple way to store small amounts of numerical data, and one that is widely used in the XAFS community, is to store data in plaintext (ASCII encoded) data files, with whitespace delimited numbers layed out as a table, with a fix number of columns and rows indicated by newlines. Typically a comment character such as "#" is used to signify header information. For instance:

```
# room temperature FeO.
# data from 20-BM, 2001, as part of NXS school
#----------------------
#   energy     xmu       i0
```

```
6911.7671  -0.35992590E-01  280101.00
6916.8730  -0.39081634E-01  278863.00
6921.7030  -0.42193483E-01  278149.00
6926.8344  -0.45165576E-01  277292.00
6931.7399  -0.47365589E-01  265707.00
```

This file and others like it can be read with the builtin `read_ascii()` function.

_io.**read_ascii**(*filename*, *comentchar='#;*%'*, *labels=None*)
   opens and read an plaintext data file, returning a new group containing the data.

   Parameters

   - **filename** (*string*) – name of file to read.

   - **commentchar** (*string*) – string of valid comment characters

   - **labels** – string to split for column labels

Some examples of `read_ascii()`:

```
larch> g = read_ascii('mydata.dat')
larch> show(g)
== Group ascii_file mydata.dat: 6 symbols ==
  attributes: <Group header attributes from mydata.dat>
  column_labels: ['energy', 'xmu', 'i0']
  energy: array<shape=(412,), type=dtype('float64')>
  filename: 'mydata.dat'
  i0: array<shape=(412,), type=dtype('float64')>
  xmu: array<shape=(412,), type=dtype('float64')>
larch>
```

which reads the data file and sets array names according to the column labels in the file. You can be explicit:

```
larch> g = read_ascii('mydata.dat', label='e mutrans monitor')
larch> show(g)
== Group ascii_file mydata.dat: 6 symbols ==
  attributes: <Group header attributes from mydata.dat>
  column_labels: ['e', 'mutrans', 'monitor']
  e: array<shape=(412,), type=dtype('float64')>
  filename: 'mydata.dat'
  monitor: array<shape=(412,), type=dtype('float64')>
  mutrans: array<shape=(412,), type=dtype('float64')>
larch>
```

and to get the data as a 2-D array:

```
larch> g  = read_ascii('mydata.dat', labels=False)
larch> show(g)
== Group ascii_file mydata.dat: 4 symbols ==
  attributes: <Group header attributes from mydata.dat>
  column_labels: []
  data: array<shape=(3, 412), type=dtype('float64')>
  filename: 'mydata.dat'
larch>
```

_io.**write_ascii**(*filename*, **args*, *commentchar='#'*, *label=None*, *header=None*)
   opens and writes arrays, scalars, and text to an ASCII file.

   Parameters

   - **commentchar** – character for comment ('#')

- **label** – array label line (autogenerated)

- **header** – array of strings for header

_io.**write_group**(*filename*, *group*, *scalars=None*, *arrays=None*, *arrays_like=None*, *commentchar='#'*)
    write data from a specified group to an ASCII data file

3.9.2 Using HDF5 Files

HDF5 is an increasingly popular data format for scientific data, as it can efficiently hold very large arrays in a heirar-chical format that holds "metadata" about the data, and can be explored with a variety of tools.

An example using h5_group() shows that one can browse through the data heirarchy of the HDF5 file, and pick out the needed data:

```
larch> g = h5group('test.h5')
larch> show(g)
== Group test.h5: 3 symbols ==
  attrs: {u'Collection Time': ': Sat Feb 4 13:29:00 2012', u'Version': '1.0.0',
          u'Beamline': 'GSECARS, 13-IDC / APS', u'Title': 'Epics Scan Data'}
  data: <Group test.h5/data>
  h5_file: <HDF5 file "test.h5" (mode r)>
larch>show(g.data)
== Group test.h5/data: 5 symbols ==
  attrs: {u'scan_prefix': '13IDC:', u'start_time': ': Sat Feb 4 13:29:00 2012',
        u'correct_deadtime': 'True', u'dimension': 2,
        u'stop_time': ': Sat Feb 4 13:44:52 2009'}
  environ: <Group test.h5/data/environ>
  full_xrf: <Group test.h5/data/full_xrf>
  merged_xrf: <Group test.h5/data/merged_xrf>
  scan: <Group test.h5/data/scan>


larch> g.data.scan.sums
<HDF5 dataset "det": shape (15, 26, 26), type "<f8">

larch> imshow(g.data.scan.sums[8:,:,:])
```

This interface is general-purpose but somewhat low-level. As HDF5 formats and schemas become standardized, better interfaces can easily be made on top of this approach.

3.10 Tutorial: Plotting and Displaying Data

Plotting and Visualizing data are vital to any scientific analysis package, and Larch provides several methods for data visualization. These are largely built on two types of data display. The first is the line plot (sometimes calld an xy plot), which shows traces of a set of functions y(x). The second type of data display supported is the 2-dimensional image display, in which a grey scale or false color map shows an image representing a 2-dimensional array of intensity.

Though not as fancy as many dedicated plotting and graphics packages, Larch attempts to provide satisfying and graphical displays of data, and the basic plots made with Larch can be high enough quality to include in publications. In addition, both line plots and image display provide interactive features such as zooming in and out, changing properties such as colors and labels. Finally, copying and saving images of the graphics is easy and can be done either with keyboard commands such as Ctrl-C or from dropdown menus on the graphic elements.

3.10.1 Line Plots

Larch provides a few functions for making line plots, with the principle function being called `plot()`. The `plot()` function takes two arrays: *x*, the abscissa array, and *y*, the ordinate array. It also accepts a very large number of optional arguments for setting properties like color, line style, labels, and so on. Most of these properties can also be set after the plot is displayed through the graphical display of the plot itself.

Multiple plot windows can be shown simultaneously and you can easily control which one to draw to.

`_plotter.`**`plot`**(*x*, *y*, ***kws*)

> **Parameters**
>
> > - **x** – array of x values
> >
> > - **y** – array of y values – same size as x
>
> Plot y(x) given 1-dimensional x and y arrays – these must be of the same size. Each x, y pair displayed is called a *trace*.
>
> There are many options for a plot, all specified by the keyword/value parameters described in the *Table of Plot Arguments* below.

`_plotter.`**`newplot`**(*x*, *y*, ***kws*)
> This is essentially the same a `plot()`, but with the option *new=True*. The rest of the arguments are as listed in *Table of Plot Arguments*.

`_plotter.`**`scatterplot`**(*x*, *y*, ***kws*)
> A scatterplot differs from a line plot in that the set of x, y values are not assumed to be in any particular order, and so are not connected with a line. Arguments are very similar to those for `plot()`, and are listed in *Table of Plot Arguments*.

`_plotter.`**`update_trace`**(*x*, *y*, *trace=1*, *win=1*, *side='left'*)
> updates an existing trace.

> **Parameters**
>
> > - **x** – array of x values
> >
> > - **y** – array of y values
> >
> > - **win** – integer index of window for plot (1 is the first window)
> >
> > - **trace** – integer index for the trace (1 is the first trace)
> >
> > - **side** – which y axis to use ('left' or 'right').

> This function is particularly useful for data to be plotted is changing and you wish to update traces from a previous `plot()` with new data without completely redrawing the entire plot. Using this method is substantially faster than replotting, and should be used for dynamic plots, such as plottting the progress of some function during a fit. Note that you cannot change properties such as color here – these will be inherited from the existing trace. In that sense, most of the properties of the trace and of the plot as a whole remain unchanged, it just happens that the data for the trace has been replaced.

Table of Plot Arguments These arguments apply for the `plot()`, `newplot()`, and `scatterplot()` methods. Except where noted, the arguments are available for `plot()` and `newplot()`. In addition, the `scatterplot()` method uses many of the same arguments for the same meaning, as indicated by the right-most column.

argument	type	default	meaning	scatterplot?
title	string	None	Plot title	yes
ylabel	string	None	abscissa label	yes
y2label	string	None	right-hand abscissa label	yes
label	string	None	trace label (defaults to 'trace N')	yes
win	integer	1	index of plot window to use (1, 2, ..., 16)	yes
side	left/right	left	side for y-axis and label	yes
grid	None/bool	None	to show grid lines	yes
color	string	blue	color to use for trace	yes
use_dates	bool	False	to show dates in xlabel (`plot()` only)	no
linewidth	int	2	linewidth for trace	no
style	string	solid	line-style for trace (solid, dashed, ...)	no
drawstyle	string	line	style connecting points of trace	no
marker	string	None	symbol to show for each point (+, o, ...)	no
markersize	int	8	size of marker shown for each point	no
dy	array	None	uncertainties for y values; error bars	no
ylog_scale	bool	False	draw y axis with log(base 10) scale	no
xmin	float	None	minimum displayed x value	yes
xmax	float	None	maximum displayed x value	yes
ymin	float	None	minimum displayed y value	yes
ymax	float	None	maximum displayed y value	yes
autoscale	bool	True	whether to automatically set plot limits	no
draw_legend	None/bool	None	whether to display legend (None: leave as is)	no
refresh	bool	True	whether to refresh display	no
	arguments that apply only for `scatterplot()`			
size	int	10	size of marker	yes
edgecolor	string	black	edge color of marker	yes
selectcolor	string	red	color for selected points	yes

For each plot window, the configuration for the plot (title, labels, grid displays, etc) and the properties of each trace (color, linewidth, ...) are preserved for the duration of that window. A few specific notes:

1. The title, label, and grid arguments to `plot()` default to `None`, which means to use the previously used value.

2. The *use_dates* option is not very rich, and simply turns x-values that are Unix timestamps into x labels showing the dates.

3. While the default is to auto-scale the plot from the data ranges, specifying any of the limits will override the corresponding limit(s).

4. The *color* argument can be any color name ("blue", "red", "black", etc), standard X11 color names ("cadetblue3", "darkgreen", etc), or an RGB hex color string of the form "#RRGGBB".

5. Valid *style* arguments are 'solid', 'dashed', 'dotted', or 'dash-dot' , with 'solid' as the default.

6. Valid *marker* arguments are '+', 'o', 'x', '^', 'v', '>', '<', '|', '_', 'square', 'diamond', 'thin diamond', 'hexagon', 'pentagon', 'tripod 1', or 'tripod 2'.

7. Valid *drawstyles* are None (which connects points with a straight line), 'steps-pre', 'steps-mid', or 'steps-post', which give a step between the points, either just after a point ('steps-pre'), midway between them ('steps-mid') or just before each point ('steps-post'). Note that if displaying discrete values as a function of time, left-to-right, and want to show a transition to a new value as a sudden step, you want
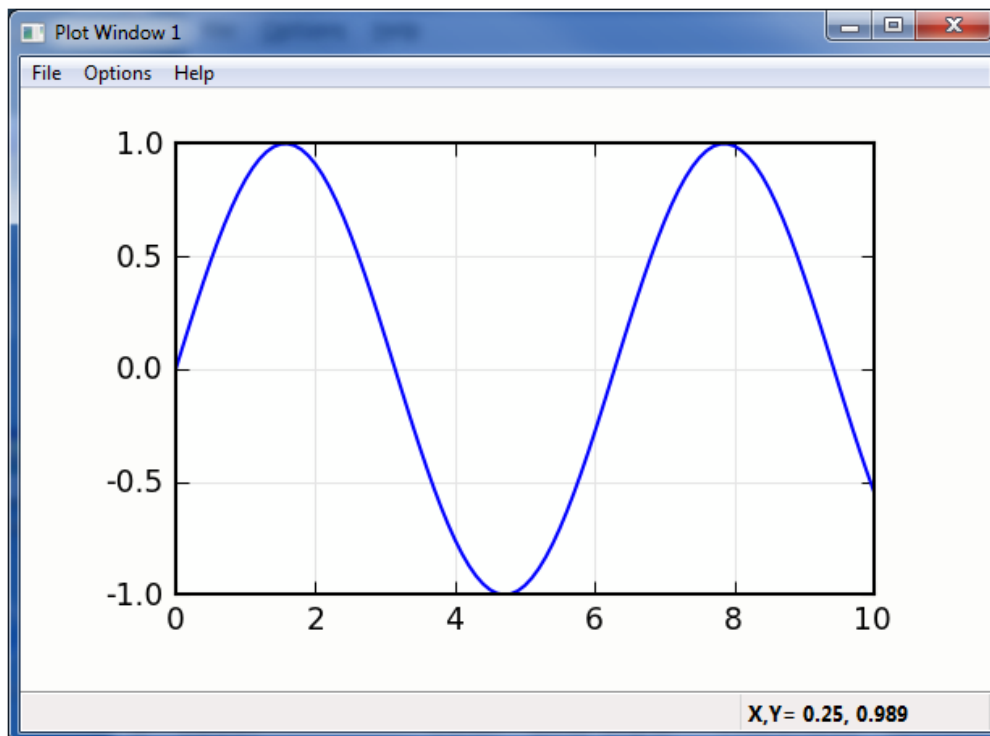
'steps-post'.

Again, most of these values can be configured interactively from the plot window.

Plot Examples

Here a a few example plots, to whet your appetite:

```
x = linspace(0, 10, 101)
y1 = sin(x)
y2 = -2 +0.2*x + (0.2*x)**2
newplot(x, y1)
```
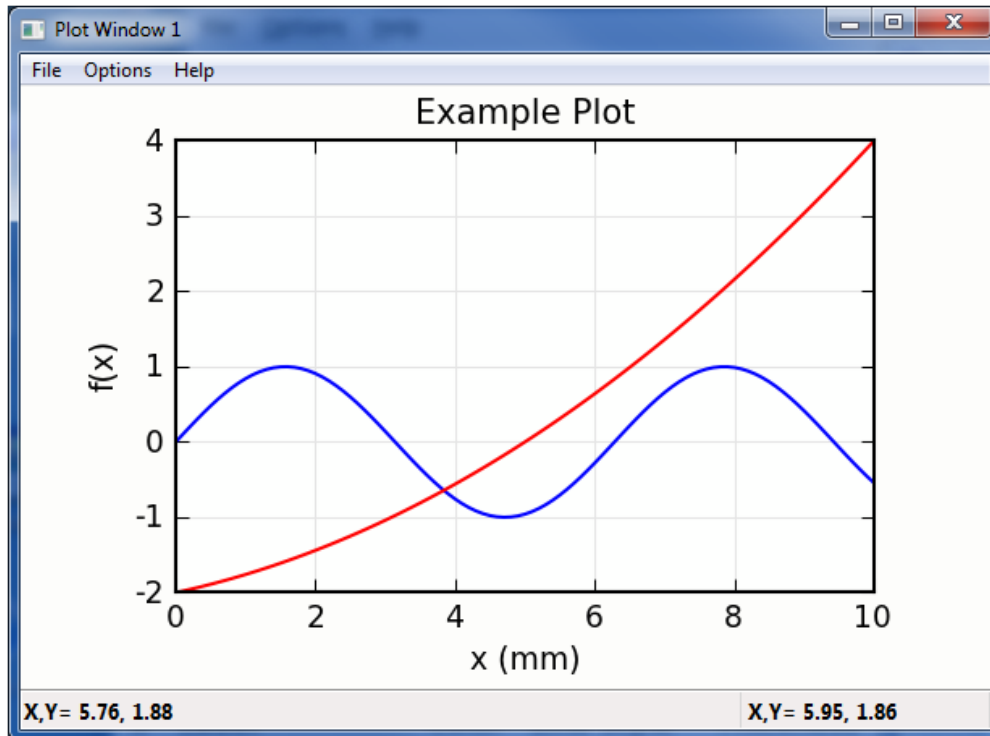
will make this plot:

Adding a second curve, and setting some labels:

```
plot(x, y2, xlabel='x (mm)', ylabel='f(x)', title='Example Plot')
```

will make this plot:

Using the Plot Windows

From the main plot window, you can perform several tasks:

Getting Cursor Position:

From the plot window you can click the left button of your mouse, and see the X, Y coordinates of where you clicked displayed in the status bar at the bottom of the plot window. You can also read the values from the variables `_plotter.plot1_x` and `_plotter.plot1_y`, for plot window 1, and `_plotter.plot2_x` and on for other plot windows.

Zooming in and out:

Left-clicking on the plot window and then dragging the mouse around with the button still pressed will draw a rectangular box around part of the plot window. Releasing the mouse will zoom in on the portion of the plot set by the rectangle. You can zoom in multiple times.

To unzoom, press Ctrl-Z (Apple-Z on Mac OS X), which will go back to the previous zoom rectangle. You can also right-click on the plot, which will bring up a window from which you can zoom out 1 level at a time, or all the way back to fully zoomed out.

Copy to Clipboard:

To copy the plot image (just the main plot image, not all the Window decorations such as menus and status bar) to the sysem clipboard, type Ctrl-C (Apple-C for Mac OS X users). You can then paste this into other applications such as rich text documents and slide presentation tools.

Save image to PNG:

To save a copy of the plot image, use Ctrl-S (Apple-S for Mac OS X users). This will bring up a 'save file' dialog box for writing a PNG file of the plot.

Print image:

On many systems, you should be able to print directly from the Plot Window, using Ctrl-P (Apple-P for Mac OS X users). This may not work on all systems.

Configuring the Plot:

From the Plot Window, either Ctrl-K (Apple-K for Mac OS X users) or Options->'Configure Plot' (or right-click to bring up a popup menu, then select Configure) will bring up the plot configuration window, which looks like this:
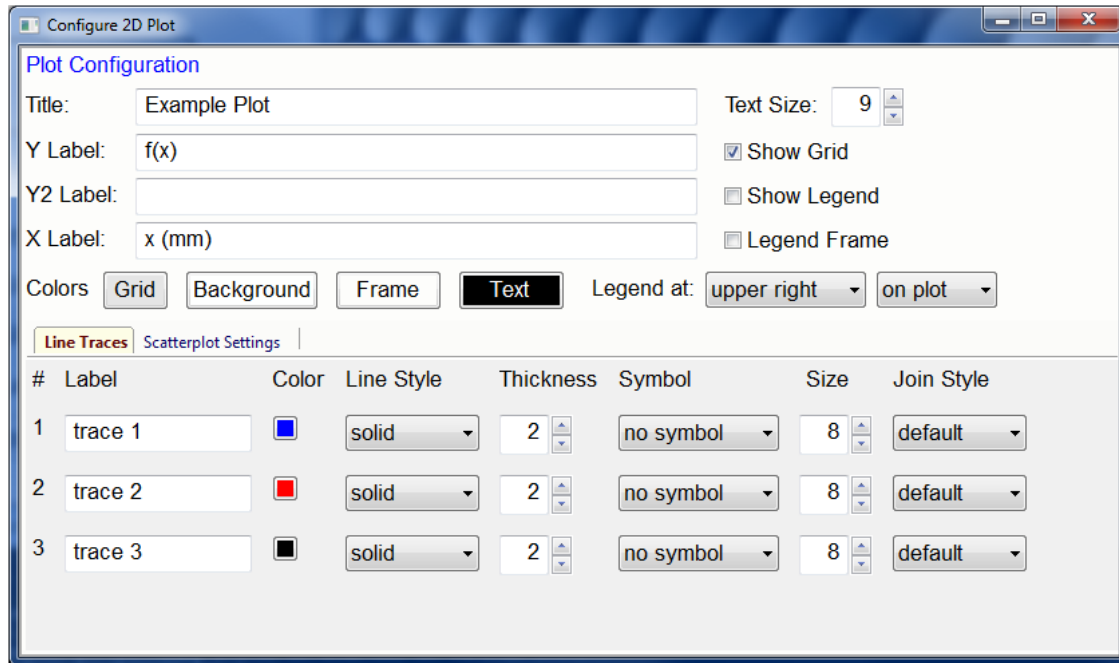
From here you can set the titles, labels, and styles, colors, symbols, and so on for the line traces.

Using TeX-like commands for labels and titles

The titles and labels for plot elements can be simple strings or use a subset of TeX markup to give fine control over typesetting greek letters, mathematical symbols and formulae. A simple example would be:

```
plot(k, chi, xlabel = ' $ k \rm(\AA^{-1}) $ ', ylabel = '$ \chi(k) $ ')
```

This will render the x and y labels as (for the pedantic, these renderings below may be only approximate):

$$k(^{-1})$$

$$\chi(k)$$

That is, strings can contain TeX-like markup between dollar signs ('$'). This does not actually use TeX (so you don't need TeX installed), and the rendering is done using the matplotlib library. For further details about using TeX for markup, including a list of symbols, commands to change fonts, and examples, can be found at http://matplotlib.sourceforge.net/users/mathtext.html

When using the Plot Configuration window to enter a TeX-like string, the text control box will be given a yellow background color (instead of the normal white color) if there is an error in rendering your TeX string.

3.10.2 Image Display

_plotter.**imshow**(*dat*, *x=None*, *y=None*, *colormap=None*, ***kws*)

> **Parameters**
>
> - **dat** – 2-d array of some intensity
> - **x** – 1-d array of x values
> - **y** – 1-d array of y values
> - **colormap** – name of color table to use.

Imshow displays a grey-scale or false-color image from a 2-d array of intensities.

The *x* and *y* arguments are intended to show real world coordinates for the image, not just the array indices. If *None*, the array indices will be shown.

By default, the image will be shown with the origin (pixel [0, 0]) in the lower left corner. The image can be rotated and flipped by the user.

The *colormap* argument is the name of the color map to use to transform intensity to color. The default table is 'gray' for a grayscale mapping. Other valid names include 'coolwarm', 'cool', 'hot', 'jet', 'Reds', 'Greens', 'Blues', 'copper', and a host of others, available

From the Image Display Window, you zoom in on regions of the image, rotate the image, change the color table, change the intensity scaling, and change the interpolattion algorithm used. An example image, generated with:

```
def gauss2d(x, y, x0, y0, sx, sy):
    return outer(exp(-(((y-y0)/float(sy))**2)/2),
                 exp(-(((x-x0)/float(sx))**2)/2))
enddef

ny, nx = 350, 400
ix = arange(nx)
iy = arange(ny)
x  =  ix / 10.
y  = -2 + iy / 10.0
dat = 0.2 + (0.5*random.random(size=nx*ny).reshape(ny, nx) +
            2.0*gauss2d(ix, iy, 190,   176,  37,  32))

imshow(dat, x=x, y=y, colormap='coolwarm')
```

is shown below:
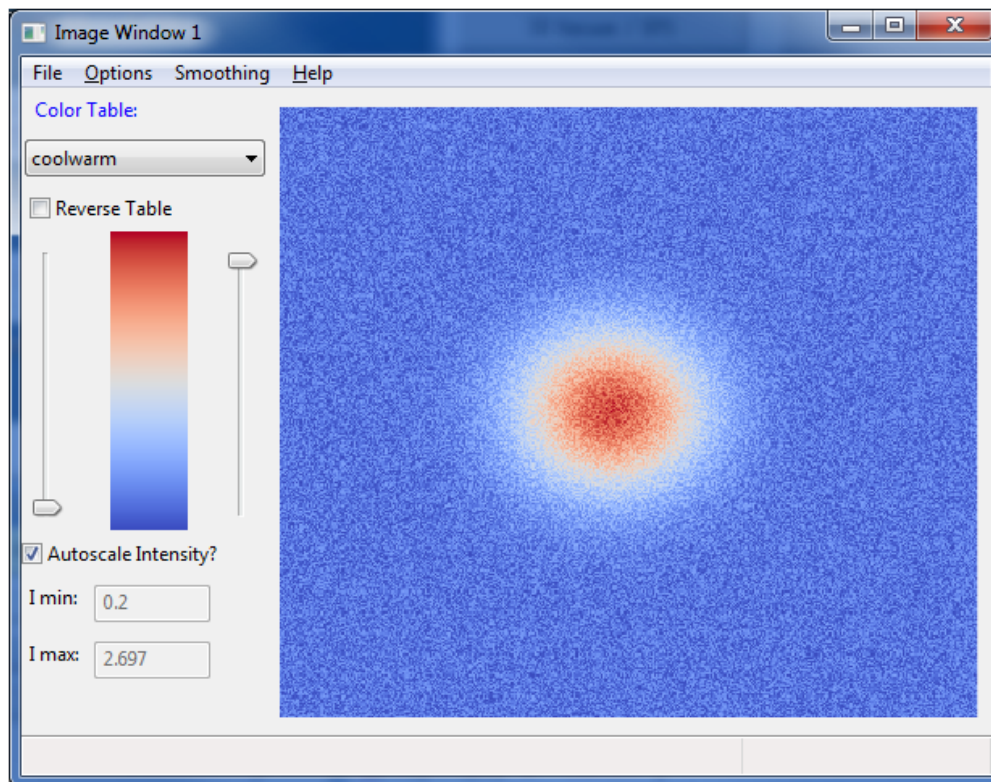
LARCH REFERENCE

This chapter describes further details of Larch language, intending to act as a reference manual. As discussed elsewhere, the single most important fact about Larch is that it implemented with and closely related to Python. Of course, Python is very well documented, and much of the Python documentation can be applied to Larch. Thus the discussion here focuses on the differences betwen Larch and Python, and on the functionality unique to Larch.

Much of the discussion here will expect a familiarity with programming and a willingness to consult the on-line Python documentation when necessary.

Needed here:

1. buit in functions

2. namespace layout (_main, _sys, _math, _builtin,...)

4.1 Larch for Developers

This chapter describes details of Larch language for developers and programmers wanting to extend Larch. This document will assume you have some familiarity with Python.

4.1.1 Differences between Larch and Python

Larch is very similar to Python but there are some very important differences, especially for someone familiar with Python. These differences are not because we feel Python is somehow inadequate or imperfect, but because Larch is a domain-specific-language. It is really something of an implementation detail that Larch's syntax is so close to Python. The principle differences with Python are:

1. Using 'end*' instead of significant white-space for code blocks.

2. Groups versus Modules

3. Changing the lookup rules for finding symbol names

4. Not implementing several python concepts, notably Class, and lambda.

Each of these is discussed in more detail below.

Some background and discussion of Larch implementation may help inform the discussion below, and help describe many of the design decisions made in Larch. First and foremost, Larch is designed to be a domain-specific macro language that makes heavy use of Python's wonderful tools for processing scientific data. Having a macro language that was similar to Python was not the primary goal. The first version of Larch actually had much weaker correspondance to Python. It turned out that the implementation was much easier and more robust when using syntax close to Python.

When larch code is run, the text is first *translated into Python code* and then parsed by Python's own *ast* module which parses Python code into an *abstract syntax tree* that is much more convenient for a machine to execute. As a brief description of what this module does, the statement:

```
a*sin(2*b)+c
```

will be parsed and translated into something like:

```
Add(Name('c'), Mult(Name('a'), Call(Name('sin'), Args([Mult(Num(2), Name('b'))]))))
```

Larch then walks through this tree, executing each Add, Mult, etc on its arguments. If you've ever done any text processing or thought about how a compiler works, you'll see that having this translation step done by proven tools is a huge benefit. For one thing, using Python's own interpreter means that Larch simply does not having parsing errors – any problem would be translation of Larch code into Python, or in executing the compiled code above. This also makes the core code implementing Larch much easier (the core functionality is fewer than 3000 lines of code).

Given this main implementation feature of Larch, you can probably see where and how the differences with Python arise:

- The Larch-to-Python translation step converts the 'end*' keywords into significant whitespace ('commenting out 'endif' etc if needed).

- The lookup for symbols in **Name('c')** is done at run-time, allowing changes from the standard Python name lookup rules.

- Unimplemented Python constructs (class, lambda, etc) are parsed, but

You can also see that Python's syntax is followed very closely, so that the translation from Larch-to-Python is minimal.

Code Block Ends

Unlike Python, Larch does not use significant whitespace to define blocks. There, that was easy. Instead, Larch uses "end blocks", of the form:

```
if test:
   <block of statements>
endif
```

Each of the Keywords *if*, *while*, *for*, *def*, and *try* must be matched with a corresponding 'end' keyword: *endif*, *endwhile*, *endfor*, *enddef*, and *endtry*. You do not need an *endelse*, *endelif*, *endexcept*, etc, as this is not ambiguous.

As a special note, you can place a '#' in front of 'end'. Note that this means exactly 1 '#' and exactly in front of 'end', so that '#endif' is allowed but not '####endif' or '# endfor'. This allows you to follow Python's indenting rules and write code that is valid Larch and valid Python, which can be useful in translating code:

```
for i in range(5)
    print(i, i/2.0)
#endfor
```

This code is both valid Larch and valid Python.

Groups vs Modules

This is at least partly a semantic distinction. Larch organizes data and code into Groups – simple containers that hold data, functions, and other groups. These are implemented as a simple, empty class that is part of the symbol table.

Symbol Lookup Rules

Looking up symbol names is a key feature of any language. Python and Larch both allow *namespaces* in which symbols can be nested into a heirarchy, using a syntax of **parent.child**, with a dot ('.') separating the components of the name. Such parent/child relationships for symbol names are used for modules (files of code), and object attributes. Thus, one could have data objects named:

```
cu_01.data.chi
cu_02.path1.chi
cu_03.model.chi
```

Where the name *chi* is used repeatedly, but with different (and multiple) parents. The issue of name lookups is how to know what (if any) to use if *chi* is specified without its parents names.

In Python, name lookups are quite straightforward and strict: "*local, module*". Here, *local* means "inside the current function or method" and *module* means "inside the current module (file of code text)". More specifically, each function or method and each module is given its own namespace, and symbols are looked for first in the local namespace, and then in the module namespace. These rules are focused on *code* rather than *data*, and leads to having a lot of "import" statements at the top of python modules. For example, to access the `sqrt()` function from the numpy module, one typically does one of these:

```python
import numpy

def sqrt_array(npts=10):
    x = np.arange(npts)/2.
    return numpy.sqrt(x)
```

or:

```python
import numpy as np

def sqrt_array(npts=10):
    x = np.arange(npts)/2.
    return np.sqrt(x)
```

In both of these examples the numpy module is brought into the *module* level namespace, either named as 'numpy' or renamed to 'np' (a common convention in scientific python code). Inside the function `sqrt_array()`, the names 'npts' and 'x' are in the local namespace – they are not available outside the function. The functions `arange()` and `sqrt()` are taken from the module-level namespace, using the name as defined in the import statement. A third alternative would be to import only the names 'sqrt' and 'arange' into the modules namespace:

```python
from numpy import sqrt, arange

def sqrt_array(npts=10):
    x = arange(npts)/2.
    return sqrt(x)
```

For quick and dirty Python scripts, there is a tendency to use "import *", as in:

```python
from numpy import *
def sqrt_array(npts=10):
    x = arange(npts)/2.
    return sqrt(x)
```

which imports several hundred names into the module level namespace. Many experienced developers will tell you to avoid this like the plague.

In Larch, the general problem of how to lookup the names of objects remains, but the rules are changed slightly. Since Group objects are used extensively throughout Larch exactly to provide namespaces as a way to organize data, we

might as well use them. Instead of using "import *", Larch has a top-level group '_math' in which it stores several hundred names of functions, mostly from the numpy module. It also uses top-level groups '_sys' and '_builtin', which hold non-mathematical builtin functions and data, and many plugins will add top-level groups (such as '_plotter', '_xafs', and '_xray'). So, to access sqrt() and arange() in Larch, you could write _math.sqrt() and _math.arange(). But you don't have to.

Symbol lookup in Larch uses a list of Groups which is searched for names. This list of groups is held in _sys.searchGroups (which holds the group names) and _sys.searchGroupObjects (which holds references to the groups themselves). These will be changed as the program runs. They can be changed dynamically, this is not encouraged (and can lead to Larch not being able to work well).

Larch also has 3 special variables that it uses to hold references to groups that are *always* included in the search of names. These are '_sys.localGroup', which holds the group for a currently running function while it is running; '_sys.moduleGroup', which holds the namespace for a module associated with a currently running function; and '_sys.paramGroup', which holds a group of Parameters used during fits (more on this, and why it is needed here in the section on Parameters).

Unimplemented features

A domain-specific-language like Larch does not need to be as full-featured as Python, so we left a few things out. These include (this may not be an exhaustive list):

- eval – Larch *is* sort of a Python eval
- lambda
- class
- global
- generators, yield
- decorators

4.1.2 Modules

Larch can import modules either written in Larch (with a '.lar' extension) or Python (with a '.py' extension). When importing a Python module, the full set of Python objects is imported as a module, which looks and acts exactly like a Group.

4.1.3 Plugins

Plugins are a powerful feature of Larch that allow it to be easily extended without the requiring detailed knowledge of all of Larch's internals. A plugin is a specially written Python module that is meant to add functionality to Larch at run-time. Generally speaking, plugins will be python modules which define new or customized versions of functions to create or manipulate Groups.

Plugins need access to Larch's symbol table and need to tell Larch how to use them. To do this, each function to be added to Larch in a plugin module needs a *_larch* keyword argument, which will be used to pass in the instance of the current larch interpreter. Normally, you will only need the *symtable* attribute of the *_larch* variable, which is the symbol table used.

In addition, all functions to be added to Larch need to be *registered*, by defining a function call registerLarchPlugin() that returns a tuple containing the name of the group containing the added functions, and a dictionary of Larch symbol names and functions. A simple plugin module would look like:

```python
def _f1(x, y, _larch=None):   # Note:  larch instance passed in to '_larch'
    if _larch is None: return
    group = _larch.symtable.create_group(name='created by f1')

    setattr(group, 'x', x) # add symbols by "setting attributes"
    setattr(group, 'y', y)

    return group

def registerLarchPlugin(): # must have a function with this name!
    return ('mymod', {'f1': _f1})
```

This is a fairly trivial example, simply putting data into a Group. Of course, the main point of a plugin is that you can do much more complicated work inside the function.

If this is placed in a file called 'myplugin.py' in the larch plugins folder (either $HOME/.larch/plugins/ or /usr/local/share/larch/plugins on Unix, or C:\Users\ME\larch\plugins or C:\Program Files\larch\plugins on Windows), then:

```
larch> add_plugin('myplugin')
```

will add a top-level group 'mymod' with an 'f1' function, so that:

```
larch> g1 = mymod.f1(10, 'yes')
larch> print g1
<Group created by f1!>
larch> print g1.x, g1.y
(10, 'yes')
```

For commonly used plugins, the `add_plugin()` call can be added to your startup script.

4.2 Overview

Larch requires Python version 2.6 or higher. Support for Python 3.X is partial, in that the core of Larch does work but is not particularly well-tested. Importantly, wxPython, the principle GUI toolkit used by Larch, has not yet been ported to Python 3.X, and so no graphical or plotting capabilities are available yet for Larch using Python 3.

LARCH FOR DEVELOPERS

This chapter describes details of Larch language for developers and programmers wanting to extend Larch. This document will assume you have some familiarity with Python.

5.1 Differences between Larch and Python

Larch is very similar to Python but there are some very important differences, especially for someone familiar with Python. These differences are not because we feel Python is somehow inadequate or imperfect, but because Larch is a domain-specific-language. It is really something of an implementation detail that Larch's syntax is so close to Python. The principle differences with Python are:

1. Using 'end*' instead of significant white-space for code blocks.

2. Groups versus Modules

3. Changing the lookup rules for finding symbol names

4. Not implementing several python concepts, notably Class, and lambda.

Each of these is discussed in more detail below.

Some background and discussion of Larch implementation may help inform the discussion below, and help describe many of the design decisions made in Larch. First and foremost, Larch is designed to be a domain-specific macro language that makes heavy use of Python's wonderful tools for processing scientific data. Having a macro language that was similar to Python was not the primary goal. The first version of Larch actually had much weaker correspondance to Python. It turned out that the implementation was much easier and more robust when using syntax close to Python.

When larch code is run, the text is first *translated into Python code* and then parsed by Python's own *ast* module which parses Python code into an *abstract syntax tree* that is much more convenient for a machine to execute. As a brief description of what this module does, the statement:

```
a*sin(2*b)+c
```

will be parsed and translated into something like:

```
Add(Name('c'), Mult(Name('a'), Call(Name('sin'), Args([Mult(Num(2), Name('b'))]))))
```

Larch then walks through this tree, executing each Add, Mult, etc on its arguments. If you've ever done any text processing or thought about how a compiler works, you'll see that having this translation step done by proven tools is a huge benefit. For one thing, using Python's own interpreter means that Larch simply does not having parsing errors – any problem would be translation of Larch code into Python, or in executing the compiled code above. This also makes the core code implementing Larch much easier (the core functionality is fewer than 3000 lines of code).

Given this main implementation feature of Larch, you can probably see where and how the differences with Python arise:

- The Larch-to-Python translation step converts the 'end*' keywords into significant whitespace ('commenting out 'endif' etc if needed).

- The lookup for symbols in **Name('c')** is done at run-time, allowing changes from the standard Python name lookup rules.

- Unimplemented Python constructs (class, lambda, etc) are parsed, but

You can also see that Python's syntax is followed very closely, so that the translation from Larch-to-Python is minimal.

5.1.1 Code Block Ends

Unlike Python, Larch does not use significant whitespace to define blocks. There, that was easy. Instead, Larch uses "end blocks", of the form:

```
if test:
    <block of statements>
endif
```

Each of the Keywords *if*, *while*, *for*, *def*, and *try* must be matched with a corresponding 'end' keyword: *endif*, *endwhile*, *endfor*, *enddef*, and *endtry*. You do not need an *endelse*, *endelif*, *endexcept*, etc, as this is not ambiguous.

As a special note, you can place a '#' in front of 'end'. Note that this means exactly 1 '#' and exactly in front of 'end', so that '#endif' is allowed but not '####endif' or '# endfor'. This allows you to follow Python's indenting rules and write code that is valid Larch and valid Python, which can be useful in translating code:

```
for i in range(5)
    print(i, i/2.0)
#endfor
```

This code is both valid Larch and valid Python.

5.1.2 Groups vs Modules

This is at least partly a semantic distinction. Larch organizes data and code into Groups – simple containers that hold data, functions, and other groups. These are implemented as a simple, empty class that is part of the symbol table.

5.1.3 Symbol Lookup Rules

Looking up symbol names is a key feature of any language. Python and Larch both allow *namespaces* in which symbols can be nested into a heirarchy, using a syntax of **parent.child**, with a dot ('.') separating the components of the name. Such parent/child relationships for symbol names are used for modules (files of code), and object attributes. Thus, one could have data objects named:

```
cu_01.data.chi
cu_02.path1.chi
cu_03.model.chi
```

Where the name *chi* is used repeatedly, but with different (and multiple) parents. The issue of name lookups is how to know what (if any) to use if *chi* is specified without its parents names.

In Python, name lookups are quite straightforward and strict: "*local, module*". Here, *local* means "inside the current function or method" and *module* means "inside the current module (file of code text)". More specifically, each function or method and each module is given its own namespace, and symbols are looked for first in the local namespace, and then in the module namespace. These rules are focused on *code* rather than *data*, and leads to having a lot of "import"

statements at the top of python modules. For example, to access the `sqrt()` function from the numpy module, one typically does one of these:

```python
import numpy

def sqrt_array(npts=10):
    x = np.arange(npts)/2.
    return numpy.sqrt(x)
```

or:

```python
import numpy as np

def sqrt_array(npts=10):
    x = np.arange(npts)/2.
    return np.sqrt(x)
```

In both of these examples the numpy module is brought into the *module* level namespace, either named as 'numpy' or renamed to 'np' (a common convention in scientific python code). Inside the function `sqrt_array()`, the names 'npts' and 'x' are in the local namespace – they are not available outside the function. The functions `arange()` and `sqrt()` are taken from the module-level namespace, using the name as defined in the import statement. A third alternative would be to import only the names 'sqrt' and 'arange' into the modules namespace:

```python
from numpy import sqrt, arange

def sqrt_array(npts=10):
    x = arange(npts)/2.
    return sqrt(x)
```

For quick and dirty Python scripts, there is a tendency to use "import *", as in:

```python
from numpy import *
def sqrt_array(npts=10):
    x = arange(npts)/2.
    return sqrt(x)
```

which imports several hundred names into the module level namespace. Many experienced developers will tell you to avoid this like the plague.

In Larch, the general problem of how to lookup the names of objects remains, but the rules are changed slightly. Since Group objects are used extensively throughout Larch exactly to provide namespaces as a way to organize data, we might as well use them. Instead of using "import *", Larch has a top-level group '_math' in which it stores several hundred names of functions, mostly from the numpy module. It also uses top-level groups '_sys' and '_builtin', which hold non-mathematical builtin functions and data, and many plugins will add top-level groups (such as '_plotter', '_xafs', and '_xray'). So, to access `sqrt()` and `arange()` in Larch, you could write *_math.sqrt()* and *_math.arange()*. But you don't have to.

Symbol lookup in Larch uses a list of Groups which is searched for names. This list of groups is held in _sys.searchGroups (which holds the group names) and _sys.searchGroupObjects (which holds references to the groups themselves). These will be changed as the program runs. They can be changed dynamically, this is not encouraged (and can lead to Larch not being able to work well).

Larch also has 3 special variables that it uses to hold references to groups that are *always* included in the search of names. These are '_sys.localGroup', which holds the group for a currently running function while it is running; '_sys.moduleGroup', which holds the namespace for a module associated with a currently running function; and '_sys.paramGroup', which holds a group of Parameters used during fits (more on this, and why it is needed here in the section on Parameters).

5.1.4 Unimplemented features

A domain-specific-language like Larch does not need to be as full-featured as Python, so we left a few things out. These include (this may not be an exhaustive list):

- eval – Larch *is* sort of a Python eval

- lambda

- class

- global

- generators, yield

- decorators

5.2 Modules

Larch can import modules either written in Larch (with a '.lar' extension) or Python (with a '.py' extension). When importing a Python module, the full set of Python objects is imported as a module, which looks and acts exactly like a Group.

5.3 Plugins

Plugins are a powerful feature of Larch that allow it to be easily extended without the requiring detailed knowledge of all of Larch's internals. A plugin is a specially written Python module that is meant to add functionality to Larch at run-time. Generally speaking, plugins will be python modules which define new or customized versions of functions to create or manipulate Groups.

Plugins need access to Larch's symbol table and need to tell Larch how to use them. To do this, each function to be added to Larch in a plugin module needs a *_larch* keyword argument, which will be used to pass in the instance of the current larch interpreter. Normally, you will only need the *symtable* attribute of the *_larch* variable, which is the symbol table used.

In addition, all functions to be added to Larch need to be *registered*, by defining a function call `registerLarchPlugin()` that returns a tuple containing the name of the group containing the added functions, and a dictionary of Larch symbol names and functions. A simple plugin module would look like:

```
def _f1(x, y, _larch=None):  # Note:  larch instance passed in to '_larch'
    if _larch is None: return
    group = _larch.symtable.create_group(name='created by f1')

    setattr(group, 'x', x) # add symbols by "setting attributes"
    setattr(group, 'y', y)

    return group

def registerLarchPlugin(): # must have a function with this name!
    return ('mymod', {'f1': _f1})
```

This is a fairly trivial example, simply putting data into a Group. Of course, the main point of a plugin is that you can do much more complicated work inside the function.

If this is placed in a file called 'myplugin.py' in the larch plugins folder (either $HOME/.larch/plugins/ or /usr/local/share/larch/plugins on Unix, or C:\Users\ME\larch\plugins or C:\Program Files\larch\plugins on Windows), then:

```
larch> add_plugin('myplugin')
```

will add a top-level group 'mymod' with an 'f1' function, so that:

```
larch> g1 = mymod.f1(10, 'yes')
larch> print g1
<Group created by f1!>
larch> print g1.x, g1.y
(10, 'yes')
```

For commonly used plugins, the `add_plugin()` call can be added to your startup script.

XAFS ANALYSIS WITH LARCH

One of the primary motivations for Larch was processing XAFS data. Larch was originally conceived to be version 2 of Ifeffit, replacing and expanding all the XAFS analysis capabilities of that package.

As of this writing (June, 2012), this replacement is approximately complete, in that most functionality of Ifeffit 1 is available in Larch. A few features of some processing steps are not fully available in Larch, and there are some slight differences in implementation details such that slightly different numerical results are obtained. On the other hand, some new features are already available with Larch that were not available with Ifeffit 1.2 and some small errors in Ifeffit 1.2 have been fixed.

XAFS Analysis can generally be broken into a few separate steps:

1. Reading in raw data.

2. Making corrections to the data, and converting to $\mu(E)$

3. Pre-edge background removal and normalization.

4. Interpreting normalized mu(E) as XANES spectra

5. Post-edge background removal, conversion to $\chi(k)$

6. XAFS Fourier Transform to $\chi(R)$

7. Reading and processing FEFF Paths from external files.

8. Fitting XAFS $\chi(k)$ to a sum of FEFF paths.

Broadly speaking, Larch can do all of these steps. The XAFS-specific functions in Larch are kept in the `_xafs` Group, which can be easily accessed, as this is in the default search path. Note that many of the XAFS functions take a **group** argument, which is a group into which resulting data are written. That is, many of the functions have several output arrays and groups. Many of of the functions will return the most fundamental result, but this will be a minimal subset of the possible outputs that would go into the supplied **group**.

The XAFS functionality is described in further detail in the sections listed below.

6.1 XAFS: Pre-edge Subtraction and Normalization

`_xafs.`**`pre_edge`**(*energy*, *mu*, *group=None*, *...*)
Pre-edge subtraction and normalization.

This performs a number of steps:

1. determine E_0 (if not supplied) from max of deriv(mu)

2. fit a line of polymonial to the region below the edge

3. fit a polymonial to the region above the edge

4. extrapolae the two curves to E_0 to determine the edge jump

Parameters

- **energy** – 1-d array of x-ray energies, in eV

- **mu** – 1-d array of $\mu(E)$

- **group** – output group

- **e0** – edge energy, E_0 in eV. If None, it will be determined here.

- **step** – edge jump. If None, it will be determined here.

- **pre1** – low E range (relative to E0) for pre-edge fit

- **pre2** – high E range (relative to E0) for pre-edge fit

- **nvict** – energy exponent to use for pre-edg fit. See Note below

- **norm1** – low E range (relative to E0) for post-edge fit

- **norm2** – high E range (relative to E0) for post-edge fit

- **nnorm** – number of terms in polynomial (that is, 1+degree) for post-edge, normalization curve. Default=3 (quadratic)

For return values, if **group** is None, the return value is (edge_step, e0).

If a **group** is supplied, return value is None, and the following data is put into the **group**:

attribute	meaning
e0	energy origin
edge_step	edge step
norm	normalized mu(E) (array)
pre_edge	pre-edge curve (array)
post_edge	post-edge, normalization curve (array)

Note: nvict gives an exponent to the energy term for the pre-edge fit. That is, a line $(mE+b)$ is fit to $\mu(E)E^{nvict}$ over the pr-edge region, E= [E0+pre1, E0+pre2].

_xafs.**find_e0** (*energy*, *mu*, *group=None*, ...)
Determine E_0, the energy threshold of the absorption edge, from the arrays energy and mu for $\mu(E)$.

This finds the point with maximum derivative with some checks to avoid spurious glitches.

Parameters

- **energy** – 1-d array of x-ray energies, in eV

- **mu** – 1-d array of mu(E)

- **group** – output group

Returns e0, the edge energy, E_0 in eV. If a group is supplied, group.e0 will also be set to this value.

6.1.1 Example

A simple example of pre-edge subtraction:

```
fname = 'fe2o3_rt1.xmu'
dat = read_ascii(fname, labels='energy xmu i0')

pre_edge(dat.energy, dat.xmu, group=dat)

show(dat)

newplot(dat.energy, dat.xmu, label=' $ \mu(E) $ ',
        xlabel='Energy (eV)',
        title='%s Pre-Edge ' % fname,
        show_legend=True)

plot(dat.energy, dat.pre_edge, label='pre-edge line',
     color='black', style='dashed' )

plot(dat.energy, dat.post_edge, label='normalization line',
     color='black', style='dotted' )
```
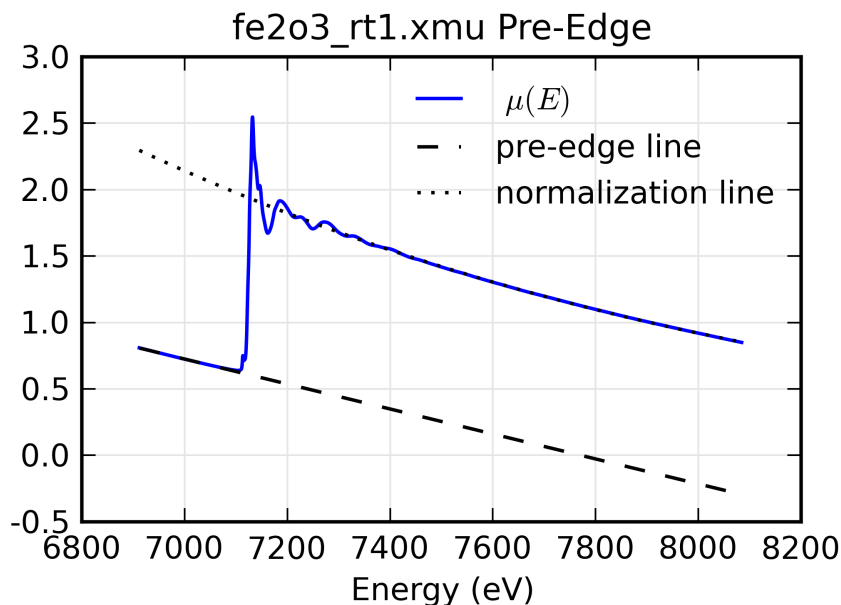
gives the following results:

6.2 XAFS: Post-edge Background Subtraction

_xafs.**autobk**(*energy*, *mu*, *group=None*, *rbkg=1.0*, *...*)
    Determine the post-edge background function, $\mu_0(E)$, and corresponding $\chi(k)$.

        Parameters

- **energy** – 1-d array of x-ray energies, in eV

- **mu** – 1-d array of $\mu(E)$

- **group** – output group

- **rbkg** – distance (in) for $\chi(R)$ above which the signal is ignored. Default = 1.

- **e0** – edge energy, in eV. If *None*, it will be determined.

- **edge_step** – edge step. If *None*, it will be determined here.

- **nknots** – number of knots in spline. If *None*, it will be determined.

- **kmin** – minimum k value [0]

- **kmax** – maximum k value [full data range].

- **kweight** – k weight for FFT. [1]

- **dk** – FFT window window parameter. [0]

- **win** – FFT window function name. ['hanning']

- **k_std** – optional k array for standard $chi(k)$.

- **chi_std** – optional χ array for standard $chi(k)$.

- **nfft** – array size to use for FFT [2048]

- **kstep** – k step size to use for FFT [0.05]

- **pre_edge_kws** – keyword arguments to pass to `pre_edge()`.

The background subtraction method used is the **AUTOBK** algorithm, in which a spline function is matched to the low-*R* components of the resulting $\chi(k)$.

For reference, $k = \sqrt{2m_e(E - E_0)/\hbar^2}$ is the wavenumber of the ejected photo-electron, where E_0 is the absorption threshold energy (the 0 of photo-electron energy). For k in units of $^{-1}$ and E in units of eV, $k \approx \sqrt{(E - E_0)/3.81}$. With this conversion of energy to wavenumber, $\chi(k)$ is defined from

$$\chi(E) = \frac{\mu(E) - \mu_0(E)}{\Delta\mu}$$

where $\mu_0(E)$ is the post-edge background function determined here, and $\Delta\mu$ is the edge step, determined from the `pre_edge()`.

6.3 XAFS: Fourier Transforms for XAFS

`_xafs.`**`ftwindow`**(*k*, *xmin=0*, *xmax=None*, *dk=1*, *…*)
 create a Fourier transform window function.

`_xafs.`**`xafsft`**(*k*, *chi*, *group=None*, *…*)
 perform a forward XAFS Fourier transform, from $\chi(k)$ to $\chi(R)$, using common XAFS conventions.

`_xafs.`**`xafsift`**(*r*, *chir*, *group=None*, *…*)
 perform an inverse XAFS Fourier transform, from $\chi(R)$ to $\chi(q)$, using common XAFS conventions.

`_xafs.`**`xafsft_fast`**(*chi*, *nfft=2048*, *kstep=0.05*)
 perform a forward XAFS Fourier transform, from $\chi(k)$ to $\chi(R)$, using common XAFS conventions. This version demands chi to represent $\chi(k)$ on a uniform k grid, and returns the complex array of $\chi(R)$ without putting any values into a group.

`_xafs.`**`xafsift_fast`**(*chi*, *nfft=2048*, *kstep=0.05*)
 perform a reverse XAFS Fourier transform, from $\chi(k)$ to $\chi(R)$, using common XAFS conventions. This version demands chir to represent the complex $\chi(R)$ as created from $\chi(k)$ on a uniform k grid, and returns the complex array of $\chi(q)$ without putting any values into a group.

6.4 XAFS: Fitting XAFS to Feff Paths

`_xafs.`**`feff_path`**(*filename...*)
> create a Feff Path object

`_xafs.`**`feffit_transform`**()
> create a Feffit Transform object

`_xafs.`**`feffit_dataset`**()
> create a Feffit Dataset object

`_xafs.`**`feffit`**()
> execute a Feffit fit.

XRF ANALYSIS WITH LARCH

X-ray Fluorescence Data can be maniuplated and displayed with Larch.

PYTHON MODULE INDEX

PYTHON MODULE INDEX

_
_io, 32
_plotter, 35
_xafs, 58

INDEX